

In [1]:

```
# IE 678 Deep Learning, University of Mannheim
# Author: Rainer Gemulla
```

In [136]:

```
# Student: Timur Michael Carstensen
# Student ID: 1722194
# Date: 27.03.2022
```

In [2]:

```
import math
import matplotlib as mpl
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F

from IPython import get_ipython
from util import nextplot
from matplotlib import inline
%matplotlib notebook
get_ipython().magic('run -i "a01-fnn-helper.py"')
```


```
/usr/local/Caskroom/miniconda/base/envs/mtp-ai-turing-tumble/lib/python3.9/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found.
Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user\_install.html (https://ipywidgets.readthedocs.io/en/stable/user\_install.html)
```

```
from .autonotebook import tqdm as notebook_tqdm
/Users/timurcarstensen/Library/CloudStorage/OneDrive-bwedu/1. Modules/1. Master/1. MMDS/2. Semester/IE 678 - Deep Learning/4-Assignments/ie-678-deep-learning/dl22-a01/a01-fnn-helper.py:26: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations (https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations)
y = np.concatenate([np.zeros(n, dtype=np.int), np.ones(n, dtype=np.int)])
```

1 Perceptrons


In [3]:

```
# plot X1 (separable)
nextplot()
plot2(X1, y1)
```



In [4]:

```
# plot X2 (not separable)  
nextplot()  
plot2(X2, y2)
```



In [48]:

```
def pt_classify(X, w):  
    """Classify using a perceptron.  
  
    Parameters  
    -----  
    X : torch array of shape (N,D) or shape (D,)   
        Design matrix of test examples  
    w : torch array of shape (D,)   
        Weight vector  
  
    Returns  
    -----  
    torch array of shape (N,)   
        Predicted binary labels (either 0 or 1)"""  
    if X.dim() == 1:  
        X = X.view(1, -1)  
    return (X @ w >= 0).int()
```

1a+c Learning

In [44]:

```

def pt_train(X, y, maxepochs=100, pocket=False, w0=None):
    """Train a perceptron.

    Parameters
    -----
    X : torch array of shape (N,D)
        Design matrix
    y : torch array of shape (N,)
        Binary labels (either 0 or 1)
    maxepochs : int
        Maximum number of passes through the training set before the algorithm
        returns
    pocket : bool
        Whether to use the pocket algorithm (True) or the perceptron learning algorithm
        (False)
    w0 : torch array of shape (D,)
        Initial weight vector

    Returns
    -----
    torch array of shape (D,)
        Fitted weight vector"""

    N, D = X.shape
    if w0 is None: # initial weight vector
        w0 = torch.zeros(D)
    w = w0 # current weight vector

    train = X.clone()
    train_target = y.clone()
    train_target[train_target==0.0] = -1

    total_weight_updates = 0
    total_tested_examples = 0
    total_correctly_classified_examples = 0

    pocket_weight_vector = w.clone()
    pocket_weight_vector_count = 0

    local_weight_vector = pocket_weight_vector.clone()
    local_weight_vector_count = 0

    for epoch in range(maxepochs):
        no_updates: bool = False

        if not pocket:
            no_updates: bool = True

            for i, x in enumerate(train):

                total_tested_examples += 1

                if torch.sign(w[1:]@x[1:]) != torch.sign(train_target[i]):
                    w[1:] += torch.sign(train_target[i]) * x[1:]
                    total_weight_updates += 1
                    no_updates = False
                else:
                    total_correctly_classified_examples += 1

```

```

elif pocket:

    for i in range(N):
        r = torch.randint(high=N, size=(1,1)).item()
        rand_sample = train[r]
        rand_sample_target = train_target[r]

        total_tested_examples += 1

        if torch.sign(local_weight_vector[1:]@rand_sample[1:]) != torch.sign(
            local_weight_vector_count = 0
            local_weight_vector[1:] += torch.sign(rand_sample_target) * rand
            total_weight_updates += 1

        elif not torch.sign(local_weight_vector[1:]@rand_sample[1:]) != torch
            local_weight_vector_count += 1
            total_correctly_classified_examples += 1

        if local_weight_vector_count >= pocket_weight_vector_count:
            pocket_weight_vector_count = local_weight_vector_count
            pocket_weight_vector = local_weight_vector.clone()

    if no_updates:
        print(f"stopped training after {epoch + 1} epochs")
        break

print(
    f"weight updates: {total_weight_updates:4}"
    f" / tested examples: {total_tested_examples:6}"
    f" / correctly classified: {total_correctly_classified_examples:6}"
    f" / incorrectly classified: {(total_tested_examples - total_correctly_class
)

if pocket:
    print(
        f"the best weight vector for the pocket algorithm classified "
        f"{pocket_weight_vector_count} samples correctly in a row"
    )
    return pocket_weight_vector
elif not pocket:
    return w

```

1b+d Experimentation


perceptron learning algorithm on separable data

In [45]:

```
# Train a perceptron using the perceptron learning algorithm and plot decision  
# boundary. You should get a perfect classification here. The decision boundary  
# should not change if you run this multiple times.  
w = pt_train(X1, y1)  
nextplot()  
plot2(X1, y1)  
plot2db(w, label="perceptron")
```

stopped training after 4 epochs

weight updates: 5 / tested examples: 800 / correctly classified:
795 / incorrectly classified: 5




perceptron learning algorithm on non-separable data

In [46]:

```
w = pt_train(X2, y2)
nextplot()
plot2(X2, y2)
plot2db(w, label="perceptron")
```

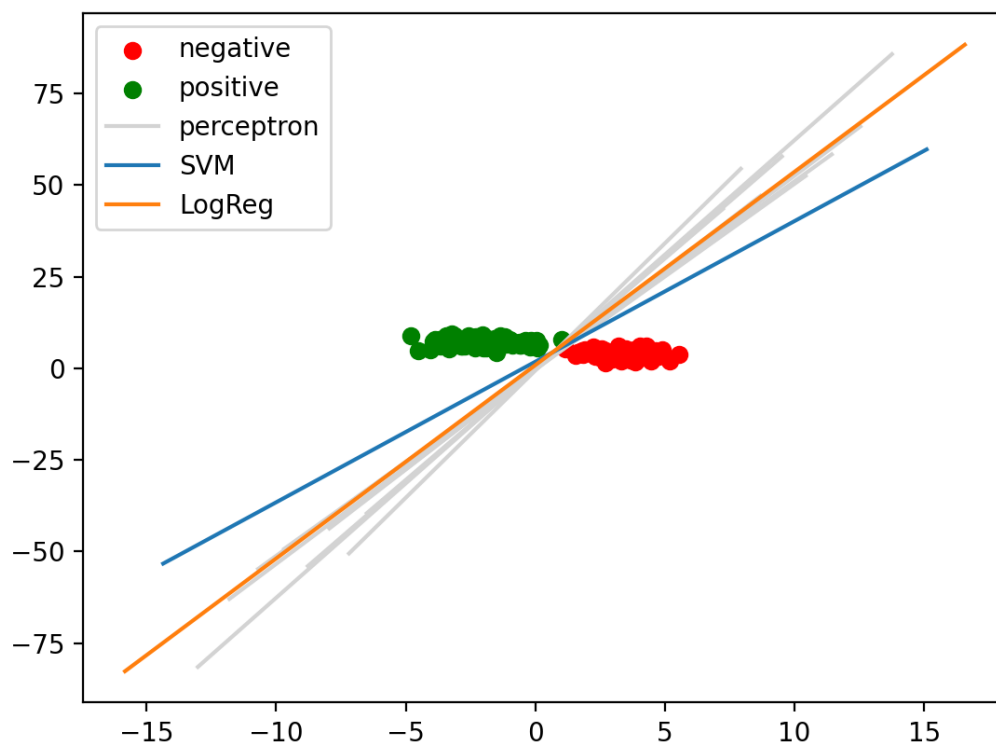
weight updates: 398 / tested examples: 20000 / correctly classified:
19602 / incorrectly classified: 398



perceptron learning algorithm on separable data compared with linear SVM and logistic regression

In [49]:

```
nextplot()
plot2dbs(X1, y1, n=10, maxepochs=1000, pocket=False)
```



```
stopped training after 4 epochs
weight updates: 5 / tested examples: 800 / correctly classified:
795 / incorrectly classified: 5
stopped training after 3 epochs
weight updates: 3 / tested examples: 600 / correctly classified:
597 / incorrectly classified: 3
stopped training after 5 epochs
weight updates: 8 / tested examples: 1000 / correctly classified:
992 / incorrectly classified: 8
stopped training after 3 epochs
weight updates: 3 / tested examples: 600 / correctly classified:
597 / incorrectly classified: 3
stopped training after 3 epochs
weight updates: 3 / tested examples: 600 / correctly classified:
597 / incorrectly classified: 3
stopped training after 5 epochs
```

```
weight updates:    8 / tested examples:    1000 / correctly classified:
992 / incorrectly classified:    8
stopped training after 6 epochs
weight updates:    10 / tested examples:    1200 / correctly classified:
1190 / incorrectly classified:    10
stopped training after 6 epochs
weight updates:    10 / tested examples:    1200 / correctly classified:
1190 / incorrectly classified:    10
stopped training after 3 epochs
weight updates:     3 / tested examples:     600 / correctly classified:
597 / incorrectly classified:     3
stopped training after 8 epochs
weight updates:    15 / tested examples:    1600 / correctly classified:
1585 / incorrectly classified:    15
```

Misclassification rates (train)

Perceptron (best result): 0


Linear SVM (C=1) : 0

Logistic regression : 0

perceptron learning algorithm on non-separable data and a comparison with linear SVM and logistic regression

In [50]:

```
nextplot()
plot2db(x2, y2, n=10, maxepochs=1000, pocket=False)
```



```
weight updates: 4127 / tested examples: 200000 / correctly classified:
195873 / incorrectly classified: 4127
weight updates: 4117 / tested examples: 200000 / correctly classified:
195883 / incorrectly classified: 4117
weight updates: 4128 / tested examples: 200000 / correctly classified:
195872 / incorrectly classified: 4128
weight updates: 4124 / tested examples: 200000 / correctly classified:
195876 / incorrectly classified: 4124
weight updates: 4122 / tested examples: 200000 / correctly classified:
195878 / incorrectly classified: 4122
weight updates: 4122 / tested examples: 200000 / correctly classified:
195878 / incorrectly classified: 4122
weight updates: 4130 / tested examples: 200000 / correctly classified:
195870 / incorrectly classified: 4130
weight updates: 4131 / tested examples: 200000 / correctly classified:
195869 / incorrectly classified: 4131
weight updates: 4132 / tested examples: 200000 / correctly classified:
195868 / incorrectly classified: 4132
weight updates: 4127 / tested examples: 200000 / correctly classified:
195873 / incorrectly classified: 4127
```


Misclassification rates (train)

```
Perceptron (best result): 1
Linear SVM (C=1)          : 3
Logistic regression       : 3
```

pocket algorithm on separable data and a comparison with linear SVM and logistic regression

In [51]:

```
nextplot()
plot2db(x1, y1, n=10, maxepochs=1000, pocket=True)
```



```
weight updates: 20 / tested examples: 200000 / correctly classified:
199980 / incorrectly classified: 20
the best weight vector for the pocket algorithm classified 197949 samp
les correctly in a row
weight updates: 5 / tested examples: 200000 / correctly classified:
199995 / incorrectly classified: 5
the best weight vector for the pocket algorithm classified 199936 samp
les correctly in a row
weight updates: 5 / tested examples: 200000 / correctly classified:
199995 / incorrectly classified: 5
the best weight vector for the pocket algorithm classified 199898 samp
les correctly in a row
weight updates: 24 / tested examples: 200000 / correctly classified:
199976 / incorrectly classified: 24
the best weight vector for the pocket algorithm classified 198814 samp
les correctly in a row
weight updates: 24 / tested examples: 200000 / correctly classified:
199976 / incorrectly classified: 24
the best weight vector for the pocket algorithm classified 197159 samp
les correctly in a row
weight updates: 5 / tested examples: 200000 / correctly classified:
199995 / incorrectly classified: 5
the best weight vector for the pocket algorithm classified 199986 samp
les correctly in a row
weight updates: 9 / tested examples: 200000 / correctly classified:
199991 / incorrectly classified: 9
the best weight vector for the pocket algorithm classified 199176 samp
les correctly in a row
weight updates: 18 / tested examples: 200000 / correctly classified:
```

```
199982 / incorrectly classified: 18
the best weight vector for the pocket algorithm classified 198033 samples correctly in a row
weight updates: 27 / tested examples: 200000 / correctly classified: 199973 / incorrectly classified: 27
the best weight vector for the pocket algorithm classified 196524 samples correctly in a row
weight updates: 14 / tested examples: 200000 / correctly classified: 199986 / incorrectly classified: 14
the best weight vector for the pocket algorithm classified 199239 samples correctly in a row
```

Misclassification rates (train)

Perceptron (best result): 0


Linear SVM (C=1) : 0

Logistic regression : 0

pocket algorithm on non-separable data and a comparison with linear SVM and logistic regression

In [53]:

```
nextplot()
plot2db(x2, y2, n=10, maxepochs=1000, pocket=True)
```



```
weight updates: 3235 / tested examples: 200000 / correctly classified:
196765 / incorrectly classified: 3235
the best weight vector for the pocket algorithm classified 1376 sample
s correctly in a row
weight updates: 3154 / tested examples: 200000 / correctly classified:
196846 / incorrectly classified: 3154
the best weight vector for the pocket algorithm classified 1493 sample
s correctly in a row
weight updates: 3232 / tested examples: 200000 / correctly classified:
196768 / incorrectly classified: 3232
the best weight vector for the pocket algorithm classified 1119 sample
s correctly in a row
weight updates: 3170 / tested examples: 200000 / correctly classified:
196830 / incorrectly classified: 3170
the best weight vector for the pocket algorithm classified 1451 sample
s correctly in a row
weight updates: 3177 / tested examples: 200000 / correctly classified:
196823 / incorrectly classified: 3177
the best weight vector for the pocket algorithm classified 1093 sample
s correctly in a row
weight updates: 3129 / tested examples: 200000 / correctly classified:
196871 / incorrectly classified: 3129
the best weight vector for the pocket algorithm classified 1532 sample
s correctly in a row
weight updates: 3172 / tested examples: 200000 / correctly classified:
196828 / incorrectly classified: 3172
the best weight vector for the pocket algorithm classified 920 samples
correctly in a row
weight updates: 2996 / tested examples: 200000 / correctly classified:
```

```
197004 / incorrectly classified: 2996
the best weight vector for the pocket algorithm classified 1221 sample
s correctly in a row
weight updates: 3229 / tested examples: 200000 / correctly classified:
196771 / incorrectly classified: 3229
the best weight vector for the pocket algorithm classified 1466 sample
s correctly in a row
weight updates: 3188 / tested examples: 200000 / correctly classified:
196812 / incorrectly classified: 3188
the best weight vector for the pocket algorithm classified 1505 sample
s correctly in a row
```


```
Misclassification rates (train)
Perceptron (best result): 1
Linear SVM (C=1)          : 3
Logistic regression       : 3
```

2 Multi-Layer Feed-Forward Neural Networks

2a Conjecture how an FNN fit will look like

In [5]:

```
# here is the one-dimensional dataset that we will use  
nextplot()  
plot1(X3, y3, label="train")  
plot1(X3test, y3test, label="test")  
plt.legend()
```



Out[5]:

```
<matplotlib.legend.Legend at 0x7f7d2b3748e0>
```

2b Train with 2 hidden units

In [37]:

```

# Training code. You do not need to modify this code.
train_bfgs = lambda model, **kwargs: train_scipy(X3, y3, model, **kwargs)

def train3(
    hidden_sizes, nreps=10, transfer=lambda: nn.Sigmoid(), train=train_bfgs, **kwargs
):
    """Train an FNN.

    hidden_sizes is a (possibly empty) list containing the sizes of the hidden layers.
    nreps refers to the number of repetitions.

    """
    best_model = None
    best_cost = math.inf
    for rep in range(nreps):
        model = fnn_model([1] + hidden_sizes + [1], transfer)
        print(f"Repetition {rep: 2d}: ", end="")
        model = train(model, **kwargs)
        mse = F.mse_loss(y3, model(X3)).item()
        if mse < best_cost:
            best_model = model
            best_cost = mse
        print(f"best_cost={best_cost:.3f}")

    return best_model

```

In [39]:

```

# Let's fit the model with one hidden layer consisting of 2 units.
model = train3([2], nreps=1)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())

```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079572

Iterations: 390

Function evaluations: 610

Gradient evaluations: 596


best_cost=0.080

Training error: 0.07957000285387039

Test error : 0.0867156982421875

In [40]:

```
# plot the data and the fit
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
# torch.linspace(0, 13, 500).unsqueeze(1): creates a tensor of dim 1 with 500
# elements that are equally spaced from 0 to 13 (i.e. unsqueeze(1) is responsible)
# for the 1d part
```



In [9]:

```
# The weight matrices and bias vectors can be read out as follows. If you want,
# use these parameters to compute the output of the network (on X3) directly and
# compare to model(X3).
for par, value in model.state_dict().items():
    print(f"{par:<15}= {value}")
```

```
linear1.weight = tensor([[ -5.0166],
                        [ 6.0141]])
linear1.bias    = tensor([13.9789, -3.4538])
output.weight   = tensor([[1.1157, 0.8249]])
output.bias     = tensor([-1.0277])
```

In [10]:

```
# now repeat this multiple times
# YOUR CODE HERE
model = train3([2], nreps=3)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573

Iterations: 364

Function evaluations: 502

Gradient evaluations: 490

best_cost=0.080

Repetition 1: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573

Iterations: 380

Function evaluations: 575

Gradient evaluations: 554

best_cost=0.080

Repetition 2: Optimization terminated successfully.

Current function value: 0.301865

Iterations: 359

Function evaluations: 410

Gradient evaluations: 410

best_cost=0.080

Training error: 0.07957139611244202

Test error : 0.08671265095472336

In [11]:

```
# From now on, always train multiple times (nreps=10 by default) and
# report best model.
model = train3([2], nreps=10)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573
Iterations: 382
Function evaluations: 596
Gradient evaluations: 587

best_cost=0.080

Repetition 1: Optimization terminated successfully.

Current function value: 0.277769
Iterations: 79
Function evaluations: 83
Gradient evaluations: 83

best_cost=0.080

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079572
Iterations: 416
Function evaluations: 577
Gradient evaluations: 566

best_cost=0.080

Repetition 3: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573
Iterations: 428
Function evaluations: 691
Gradient evaluations: 673

best_cost=0.080

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573
Iterations: 362
Function evaluations: 490
Gradient evaluations: 479

best_cost=0.080

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.286909
Iterations: 326
Function evaluations: 444
Gradient evaluations: 435

best_cost=0.080

Repetition 6: Optimization terminated successfully.

Current function value: 0.357250
Iterations: 62
Function evaluations: 70
Gradient evaluations: 70

best_cost=0.080

Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573
Iterations: 366
Function evaluations: 583
Gradient evaluations: 571

```

best_cost=0.080
Repetition 8: Warning: Desired error not necessarily achieved due to
precision loss.
    Current function value: 0.079573
    Iterations: 394
    Function evaluations: 605
    Gradient evaluations: 593
best_cost=0.080
Repetition 9: Warning: Desired error not necessarily achieved due to
precision loss.
    Current function value: 0.079573
    Iterations: 391
    Function evaluations: 535
    Gradient evaluations: 525
best_cost=0.080
Training error: 0.07956835627555847
Test error      : 0.08670976012945175

```

2c Width

In [12]:

```

# Experiment with different hidden layer sizes. To avoid recomputing
# models, you may want to save your models using torch.save(model, filename) and
# load them again using torch.load(filename).
# model1 = train3([1], nreps=10)
# model2 = train3([2], nreps=10)
# model3 = train3([3], nreps=10)
# model10 = train3([10], nreps=10)
# model50 = train3([50], nreps=10)
# model100 = train3([100], nreps=10)

```

saving models

In []:

```

torch.save(model1, "models/model1.txt")
torch.save(model2, "models/model2.txt")
torch.save(model3, "models/model3.txt")
torch.save(model10, "models/model10.txt")
torch.save(model50, "models/model50.txt")
torch.save(model100, "models/model100.txt")

```

loading models

In [14]:

```


model1 = torch.load("models/model1.txt")
model2 = torch.load("models/model2.txt")
model3 = torch.load("models/model3.txt")
model10 = torch.load("models/model10.txt")
model50 = torch.load("models/model50.txt")
model100 = torch.load("models/model100.txt")

```

1 hidden unit

In [15]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model1)
plt.title("1 hidden unit")
print("Training error:", F.mse_loss(y3, model1(X3)).item())
print("Test error      :", F.mse_loss(y3test, model1(X3test)).item())
```




Training error: 0.37291884422302246

Test error : 0.37431666254997253

2 hidden units

In [16]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model2)
plt.title("2 hidden units")
print("Training error:", F.mse_loss(y3, model2(X3)).item())
print("Test error    :", F.mse_loss(y3test, model2(X3test)).item())
```




Training error: 0.07957205921411514
Test error : 0.08671297132968903

3 hidden units

In [17]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model3)
plt.title("3 hidden units")
print("Training error:", F.mse_loss(y3, model3(X3)).item())
print("Test error      :", F.mse_loss(y3test, model3(X3test)).item())
```




Training error: 0.007324173580855131
Test error : 0.010335267521440983

10 hidden units

In [25]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model10)
plt.title("10 hidden units")
print("Training error:", F.mse_loss(y3, model10(X3)).item())
print("Test error    :", F.mse_loss(y3test, model10(X3test)).item())
```

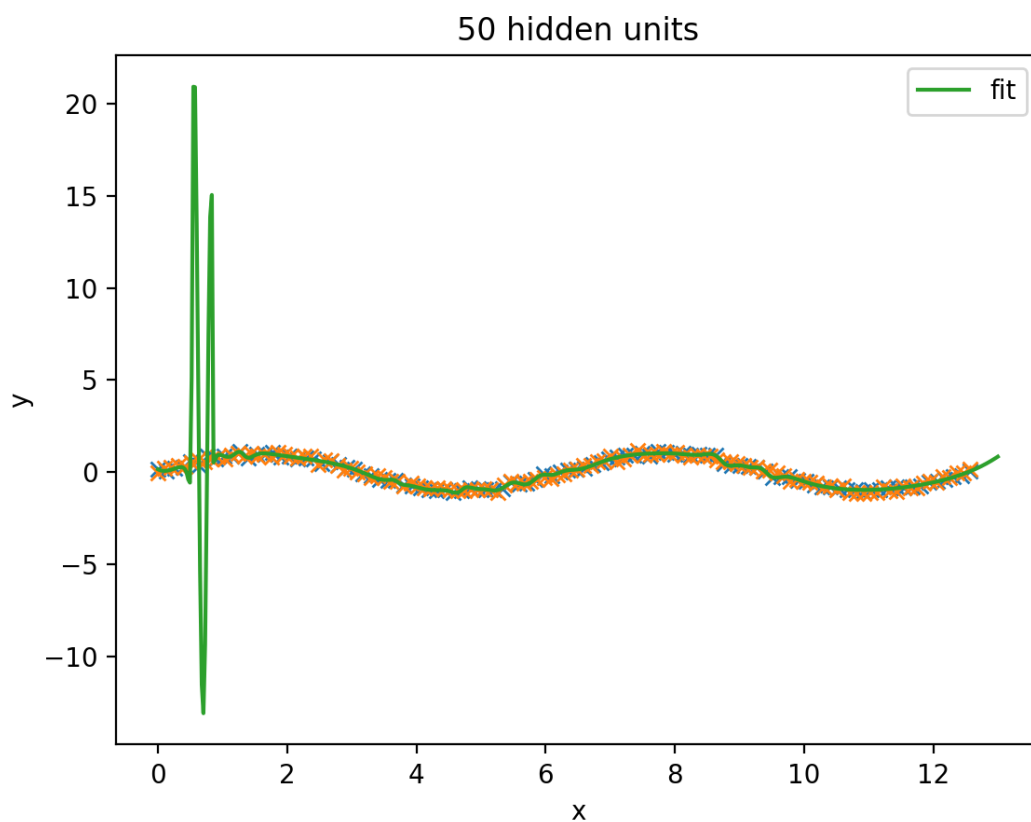


Training error: 0.006357043981552124
Test error : 0.011850706301629543

50 hidden units

In [19]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model50)
plt.title("50 hidden units")
print("Training error:", F.mse_loss(y3, model50(X3)).item())
print("Test error      :", F.mse_loss(y3test, model50(X3test)).item())
```



Training error: 0.0016042940551415086

Test error : 3.343168020248413

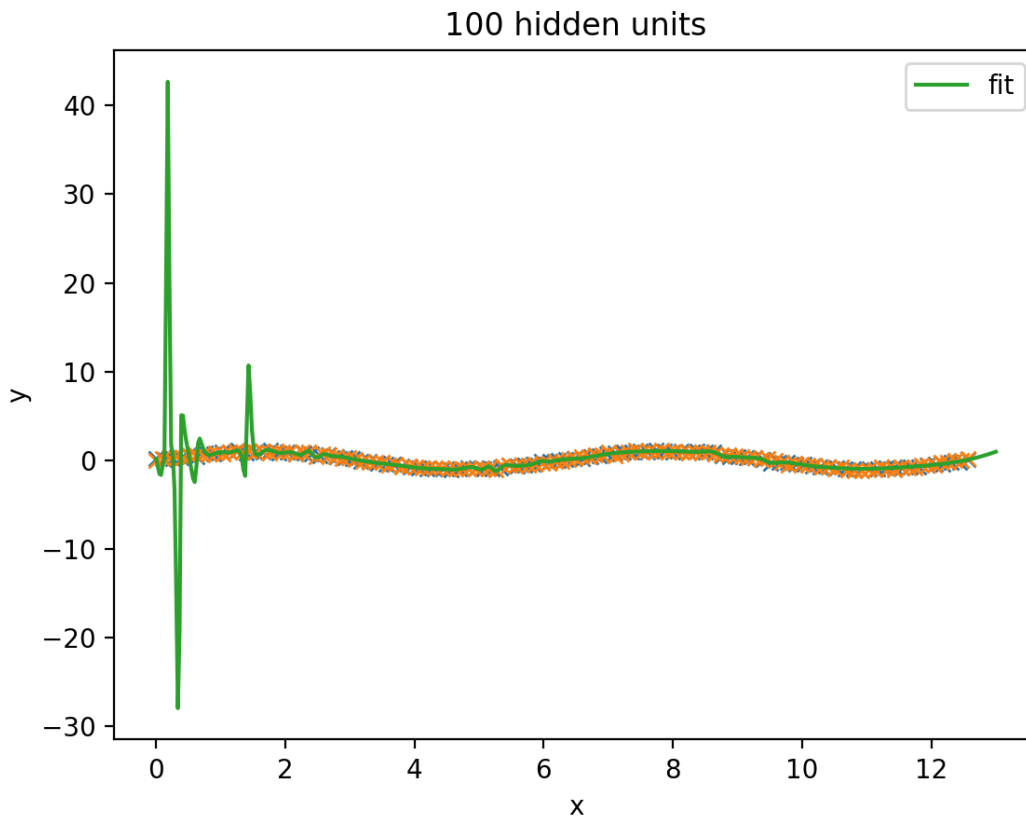
100 hidden units

In [26]:

```

nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model100)
plt.title("100 hidden units")
print("Training error:", F.mse_loss(y3, model100(X3)).item())
print("Test error      :", F.mse_loss(y3test, model100(X3test)).item())

```



Training error: 0.0010693169897422194

Test error : 6.358546733856201

Then plot the dataset as well as the predictions of each FNN on the test set into a single plot.


In [30]:

```

nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model1, label="1")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model2, label="2")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model3, label="3")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model10, label="10")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model50, label="50")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model100, label="100")
#plt.title("100 hidden units")

plt.legend(prop={'size': 13})
plt.show()

```



Training and test error for 1 to 100 hidden neurons

In [22]:

```
train1 = F.mse_loss(y3, model1(X3)).item()
test1 = F.mse_loss(y3test, model1(X3test)).item()

train2 = F.mse_loss(y3, model2(X3)).item()
test2 = F.mse_loss(y3test, model2(X3test)).item()

train3 = F.mse_loss(y3, model3(X3)).item()
test3 = F.mse_loss(y3test, model3(X3test)).item()

train10 = F.mse_loss(y3, model10(X3)).item()
test10 = F.mse_loss(y3test, model10(X3test)).item()

train50 = F.mse_loss(y3, model50(X3)).item()
test50 = F.mse_loss(y3test, model50(X3test)).item()

train100 = F.mse_loss(y3, model100(X3)).item()
test100 = F.mse_loss(y3test, model100(X3test)).item()
```

In [32]:

```
labels = ['1', '2', '3', '10', '50', '100']
train = [train1, train2, train3, train10, train50, train100]
test = [test1, test2, test3, test10, test50, test100]

x = np.arange(len(labels)) # the label locations
width = 0.4 # the width of the bars


fig, ax = plt.subplots()
rects1 = ax.bar(x - width/2, train, width, label='Train')
rects2 = ax.bar(x + width/2, test, width, label='Test')

# Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_ylabel('MSE')
ax.set_xticks(x, labels)
ax.legend(prop={'size': 15})

ax.bar_label(rects1, padding=3, fmt="%1.3f")
ax.bar_label(rects2, padding=3, fmt="%1.3f")

fig.tight_layout()

plt.show()
```



2d Distributed representations

2 hidden unit

In [63]:

```
# train a model to analyze
model_dist_2 = train3([2])
torch.save(model_dist_2, "models/model_dist_2.txt")
model_dist_2 = torch.load("models/model_dist_2.txt")
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.367556

Iterations: 54

Function evaluations: 130

Gradient evaluations: 118

best_cost=0.368

Repetition 1: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.286909

Iterations: 379

Function evaluations: 513

Gradient evaluations: 501

best_cost=0.287

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079572

Iterations: 414

Function evaluations: 589

Gradient evaluations: 575

best_cost=0.080

Repetition 3: Optimization terminated successfully.

Current function value: 0.357250

Iterations: 86

Function evaluations: 96

Gradient evaluations: 96

best_cost=0.080

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573

Iterations: 387

Function evaluations: 553

Gradient evaluations: 542

best_cost=0.080

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573

Iterations: 369

Function evaluations: 523

Gradient evaluations: 513

best_cost=0.080

Repetition 6: Optimization terminated successfully.

Current function value: 0.302737

Iterations: 127

Function evaluations: 151

Gradient evaluations: 151

best_cost=0.080

Repetition 7: Optimization terminated successfully.

Current function value: 0.357250

Iterations: 86

Function evaluations: 93

Gradient evaluations: 93

best_cost=0.080

Repetition 8: Optimization terminated successfully.

Current function value: 0.357250

Iterations: 87

Function evaluations: 99

Gradient evaluations: 99

best_cost=0.080

Repetition 9: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.079573

Iterations: 383


Function evaluations: 596

Gradient evaluations: 582

best_cost=0.080


In [87]:

```
# plot the fit as well as the outputs of each neuron in the hidden
# layer (scale for the latter is shown on right y-axis)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_2, hidden=True, scale=F
```



In [88]:

```
# plot the fit as well as the outputs of each neuron in the hidden layer, scaled
# by its weight for the output neuron (scale for the latter is shown on right
# y-axis)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_2, hidden=True, scale=1
```



weights of the model with two (2) hidden units

In [67]:

```
for par, value in model_dist_2.state_dict().items():
    print(f"{par:<15}= {value}")
```

```
linear1.weight = tensor([[ -1.0684],
                        [-1.0569]])
linear1.bias    = tensor([ 6.7106,  6.6383])
output.weight   = tensor([[ -644.4341,  646.1637]])
output.bias     = tensor([-0.8657])
```


3 hidden units

In [57]:

```
# model_dist_3 = train3([3])
# torch.save(model_dist_3, "models/model_dist_3.txt")
model_dist_3 = torch.load("models/model_dist_3.txt")
```

In [89]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_3, hidden=True, scale=F
```




In [92]:

```

nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_3, hidden=True, scale=1

```



weights of the model with three (3) hidden units

In [68]:

```

for par, value in model_dist_3.state_dict().items():
    print(f"{par:<15}= {value}")

linear1.weight = tensor([[ 0.8628],
                        [-0.8973],
                        [-0.2218]])
linear1.bias    = tensor([-8.0104,  2.9424,  1.4495])
output.weight   = tensor([[ -15.7720,  14.7583, -47.9800]])
output.bias     = tensor([24.8621])

```


10 hidden units

In [60]:

```
# model_dist_10 = train3([10])
# torch.save(model_dist_10, "models/model_dist_10.txt")
model_dist_10 = torch.load("models/model_dist_10.txt")
```

In [94]:

```
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_10, hidden=True, scale=
```

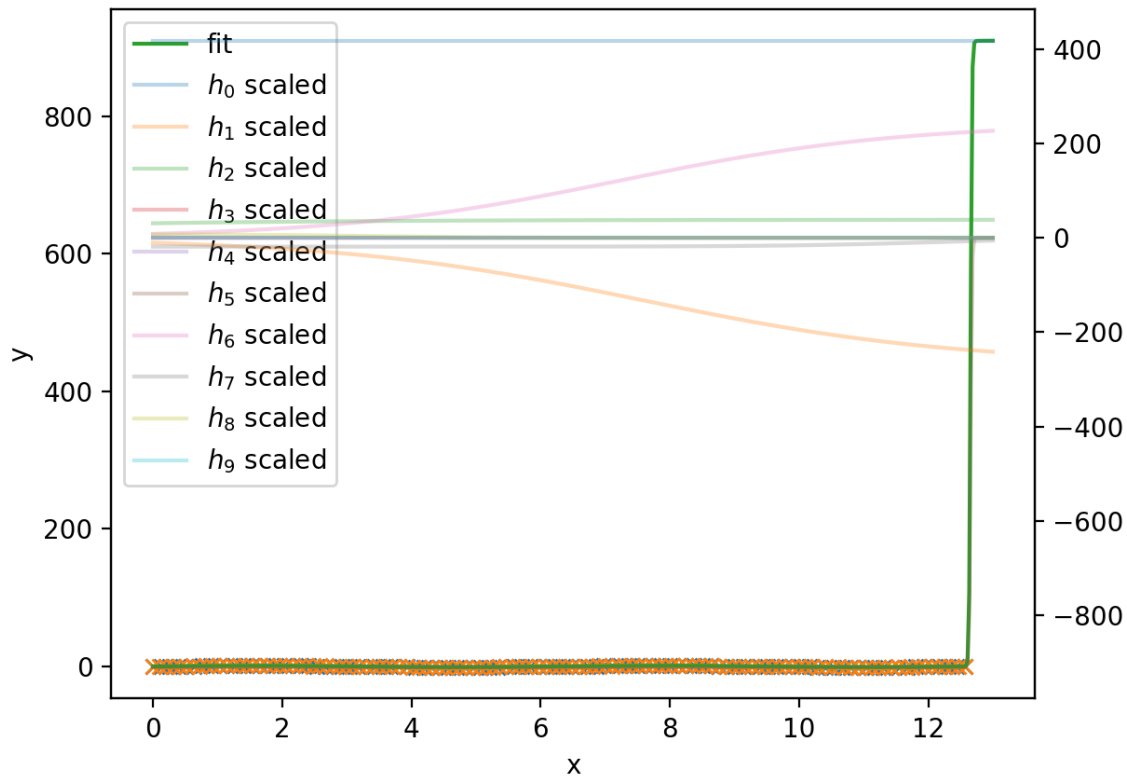


In [95]:

```

nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model_dist_10, hidden=True, scale=

```



weights of the model with ten (10) hidden units

In [70]:

```
for par, value in model_dist_10.state_dict().items():  
    print(f"{par:<15}= {value}")
```

```
linear1.weight = tensor([[ 2.8365e+02,  
    [ 4.1877e-01],  
    [ 3.3326e-01],  
    [-3.7034e+02],  
    [ 6.7506e+01],  
    [-9.8656e+01],  
    [ 4.6061e-01],  
    [-8.3900e-01],  
    [-1.0473e+00],  
    [-7.8167e+01]])  
linear1.bias   = tensor([ 365.0792,   -3.1908,    1.4124,   -7.4229, -  
587.2411, 1248.5552,  
    -3.3290,   10.1059,    3.4771, -219.2267])  
output.weight  = tensor([[ 4.1737e+02, -2.6671e+02,  3.8250e+01,  4.37  
23e+02, -2.6638e-01,  
    -9.0956e+02,  2.4256e+02, -1.8766e+01,  7.4181e+00, -2.7525e+  
01]])  
output.bias    = tensor([475.0463])
```

2e Experiment with different optimizers (optional)

Varying layers

In [129]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01), **kwargs
)
model = train3([2, 2], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=0)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.048
Repetition 1: best_cost=0.048
Repetition 2: best_cost=0.048
Repetition 3: best_cost=0.048
Repetition 4: best_cost=0.048
Repetition 5: best_cost=0.048
Repetition 6: best_cost=0.048
Repetition 7: best_cost=0.048
Repetition 8: best_cost=0.048
Repetition 9: best_cost=0.048

```



In [130]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01, **kwargs
)
model = train3([2, 2, 2], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, ver
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.373
Repetition 1: best_cost=0.277
Repetition 2: best_cost=0.052
Repetition 3: best_cost=0.052
Repetition 4: best_cost=0.052
Repetition 5: best_cost=0.052
Repetition 6: best_cost=0.028
Repetition 7: best_cost=0.028
Repetition 8: best_cost=0.028
Repetition 9: best_cost=0.028

```



In [131]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01, **kwargs
)
model = train3([3, 1, 3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, ver
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.373
Repetition 1: best_cost=0.373
Repetition 2: best_cost=0.373
Repetition 3: best_cost=0.373
Repetition 4: best_cost=0.373
Repetition 5: best_cost=0.373
Repetition 6: best_cost=0.373
Repetition 7: best_cost=0.373
Repetition 8: best_cost=0.373
Repetition 9: best_cost=0.373

```



In [135]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01), **kwargs
)
model = train3([3, 2, 2], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, ver
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.260
Repetition 1: best_cost=0.260
Repetition 2: best_cost=0.260
Repetition 3: best_cost=0.260
Repetition 4: best_cost=0.260
Repetition 5: best_cost=0.025
Repetition 6: best_cost=0.025
Repetition 7: best_cost=0.025
Repetition 8: best_cost=0.025
Repetition 9: best_cost=0.025

```



In [133]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01), **kwargs
)
model = train3([3, 3, 3, 3, 3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.373
Repetition 1: best_cost=0.041
Repetition 2: best_cost=0.041
Repetition 3: best_cost=0.041
Repetition 4: best_cost=0.041
Repetition 5: best_cost=0.041
Repetition 6: best_cost=0.041
Repetition 7: best_cost=0.041
Repetition 8: best_cost=0.041
Repetition 9: best_cost=0.041

```



Adam

In [123]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adam(model.parameters()), lr=0.01), **kwargs
)
model = train3([3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=False)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.018
Repetition 1: best_cost=0.018
Repetition 2: best_cost=0.018
Repetition 3: best_cost=0.018
Repetition 4: best_cost=0.018
Repetition 5: best_cost=0.018
Repetition 6: best_cost=0.018
Repetition 7: best_cost=0.018
Repetition 8: best_cost=0.018
Repetition 9: best_cost=0.018

```



LBFGS

In [124]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.LBFGS(model.parameters(), lr=0.01), **kwargs
)
model = train3([3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=False)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.357
Repetition 1: best_cost=0.007
Repetition 2: best_cost=0.007
Repetition 3: best_cost=0.007
Repetition 4: best_cost=0.007
Repetition 5: best_cost=0.007
Repetition 6: best_cost=0.007
Repetition 7: best_cost=0.007
Repetition 8: best_cost=0.007
Repetition 9: best_cost=0.007

```



RMSprop

In [125]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.RMSprop(model.parameters(), lr=0.01), **kwargs
)
model = train3([3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=False)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.065
Repetition 1: best_cost=0.065
Repetition 2: best_cost=0.050
Repetition 3: best_cost=0.049
Repetition 4: best_cost=0.049
Repetition 5: best_cost=0.049
Repetition 6: best_cost=0.049
Repetition 7: best_cost=0.049
Repetition 8: best_cost=0.049
Repetition 9: best_cost=0.049

```



Adagrad

In [126]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adagrad(model.parameters(), lr=0.01), **kwargs
)
model = train3([3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=False)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.440
Repetition 1: best_cost=0.426
Repetition 2: best_cost=0.426
Repetition 3: best_cost=0.419
Repetition 4: best_cost=0.412
Repetition 5: best_cost=0.412
Repetition 6: best_cost=0.412
Repetition 7: best_cost=0.412
Repetition 8: best_cost=0.412
Repetition 9: best_cost=0.412

```



Adadelta

In [127]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.Adadelta(model.parameters(), lr=0.01), **kwargs
)
model = train3([3], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, verbose=False)
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.418
Repetition 1: best_cost=0.418
Repetition 2: best_cost=0.418
Repetition 3: best_cost=0.418
Repetition 4: best_cost=0.418
Repetition 5: best_cost=0.418
Repetition 6: best_cost=0.418
Repetition 7: best_cost=0.418
Repetition 8: best_cost=0.418
Repetition 9: best_cost=0.418

```



SGD

In [128]:

```


# PyTorch provides many gradient-based optimizers; see
# https://pytorch.org/docs/stable/optim.html. You can use a PyTorch optimizer
# as follows.
train_adam = lambda model, **kwargs: fnn_train(
    X3, y3, model, optimizer=torch.optim.SGD(model.parameters(), lr=0.01), **kwargs
)
model = train3([2, 3, 2], nreps=10, train=train_adam, max_epochs=5000, tol=1e-8, ver
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)

```

```

Repetition 0: best_cost=0.506
Repetition 1: best_cost=0.506
Repetition 2: best_cost=0.505
Repetition 3: best_cost=0.503
Repetition 4: best_cost=0.488
Repetition 5: best_cost=0.488
Repetition 6: best_cost=0.488
Repetition 7: best_cost=0.488
Repetition 8: best_cost=0.488
Repetition 9: best_cost=0.488

```



In []:

```
# Experiment with different number of layers and activation functions. Here is
# an example with three hidden layers (of sizes 4, 5, and 6) and ReLU activations.
#
# You can also plot the outputs of the hidden neurons in the first layer (using
# the same code above).
model = train3([5,10,50], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Trying 1 hidden layer with a varying number of ReLUs

In [111]:

```
model = train3([1], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 1

Function evaluations: 3

Gradient evaluations: 3

best_cost=0.506

Repetition 1: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 2

Function evaluations: 3

Gradient evaluations: 3

best_cost=0.506

Repetition 2: Optimization terminated successfully.

Current function value: 0.438543

Iterations: 10

Function evaluations: 16

Gradient evaluations: 16

best_cost=0.439

Repetition 3: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 10

Function evaluations: 12

Gradient evaluations: 12

best_cost=0.439

Repetition 4: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 2

Function evaluations: 4

Gradient evaluations: 4

best_cost=0.439

Repetition 5: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 2

Function evaluations: 4

Gradient evaluations: 4

best_cost=0.439

Repetition 6: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 2

Function evaluations: 4

Gradient evaluations: 4

best_cost=0.439

Repetition 7: Optimization terminated successfully.

Current function value: 0.505991

Iterations: 24

Function evaluations: 27

Gradient evaluations: 27


best_cost=0.439

Repetition 8: Optimization terminated successfully.

Current function value: 0.506238

Iterations: 1

```
Function evaluations: 3
Gradient evaluations: 3
best_cost=0.439
Repetition 9: Optimization terminated successfully.
Current function value: 0.438543
Iterations: 8
Function evaluations: 10
Gradient evaluations: 10
best_cost=0.439
```



```
Training error: 0.4385433495044708
Test error    : 0.4407091736793518
```

In [112]:

```

model = train3([2], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())

```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406719
Iterations: 22
Function evaluations: 120
Gradient evaluations: 112

best_cost=0.407

Repetition 1: Optimization terminated successfully.

Current function value: 0.506238
Iterations: 8
Function evaluations: 9
Gradient evaluations: 9

best_cost=0.407

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406719
Iterations: 17
Function evaluations: 130
Gradient evaluations: 121

best_cost=0.407

Repetition 3: Optimization terminated successfully.

Current function value: 0.506238
Iterations: 2
Function evaluations: 4
Gradient evaluations: 4

best_cost=0.407

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406723
Iterations: 10
Function evaluations: 85
Gradient evaluations: 79

best_cost=0.407

Repetition 5: Optimization terminated successfully.

Current function value: 0.435275
Iterations: 49
Function evaluations: 67
Gradient evaluations: 67

best_cost=0.407

Repetition 6: Warning: Desired error not necessarily achieved due to precision loss.


Current function value: 0.504720
Iterations: 5
Function evaluations: 96
Gradient evaluations: 92

best_cost=0.407

Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406720
Iterations: 20
Function evaluations: 98

```
Gradient evaluations: 90
best_cost=0.407
Repetition 8: Optimization terminated successfully.
Current function value: 0.437668
Iterations: 17
Function evaluations: 19
Gradient evaluations: 19
best_cost=0.407
Repetition 9: Optimization terminated successfully.
Current function value: 0.506238
Iterations: 3
Function evaluations: 5
Gradient evaluations: 5
best_cost=0.407
```



```
Training error: 0.4067193567752838
Test error      : 0.40883004665374756
```


In [113]:

```
model = train3([3], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406722

Iterations: 13

Function evaluations: 101

Gradient evaluations: 94

best_cost=0.407

Repetition 1: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.437669

Iterations: 23

Function evaluations: 112

Gradient evaluations: 106

best_cost=0.407

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.357718

Iterations: 27

Function evaluations: 118

Gradient evaluations: 115

best_cost=0.358

Repetition 3: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406719

Iterations: 18

Function evaluations: 94

Gradient evaluations: 85

best_cost=0.358

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406720

Iterations: 16

Function evaluations: 92

Gradient evaluations: 85

best_cost=0.358

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406723

Iterations: 14

Function evaluations: 100

Gradient evaluations: 95

best_cost=0.358

Repetition 6: Optimization terminated successfully.

Current function value: 0.357620

Iterations: 34

Function evaluations: 40


Gradient evaluations: 40

best_cost=0.358

Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.406722

```
Iterations: 30
Function evaluations: 120
Gradient evaluations: 112
best_cost=0.358
Repetition 8: Warning: Desired error not necessarily achieved due to
precision loss.
Current function value: 0.406735
Iterations: 11
Function evaluations: 84
Gradient evaluations: 76
best_cost=0.358
Repetition 9: Warning: Desired error not necessarily achieved due to
precision loss.
Current function value: 0.406724
Iterations: 20
Function evaluations: 98
Gradient evaluations: 96
best_cost=0.358
```



```
Training error: 0.35762014985084534
Test error      : 0.360356867313385
```

In [114]:

```
model = train3([10], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.309325

Iterations: 35

Function evaluations: 117

Gradient evaluations: 108

best_cost=0.309

Repetition 1: Optimization terminated successfully.

Current function value: 0.085044

Iterations: 107

Function evaluations: 197

Gradient evaluations: 196

best_cost=0.085

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.084857

Iterations: 145

Function evaluations: 271

Gradient evaluations: 267

best_cost=0.085

Repetition 3: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.082022

Iterations: 110

Function evaluations: 216

Gradient evaluations: 207

best_cost=0.082

Repetition 4: Optimization terminated successfully.

Current function value: 0.356643

Iterations: 68

Function evaluations: 80

Gradient evaluations: 80

best_cost=0.082

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.013686

Iterations: 89

Function evaluations: 165

Gradient evaluations: 159

best_cost=0.014

Repetition 6: Optimization terminated successfully.

Current function value: 0.356162

Iterations: 47

Function evaluations: 56

Gradient evaluations: 56

best_cost=0.014


Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.356805

Iterations: 29

Function evaluations: 121

```
Gradient evaluations: 112
best_cost=0.014
Repetition 8: Warning: Desired error not necessarily achieved due to
precision loss.
Current function value: 0.081694
Iterations: 111
Function evaluations: 204
Gradient evaluations: 197
best_cost=0.014
Repetition 9: Warning: Desired error not necessarily achieved due to
precision loss.
Current function value: 0.356233
Iterations: 32
Function evaluations: 121
Gradient evaluations: 115
best_cost=0.014
```



```
Training error: 0.013686132617294788
Test error      : 0.017258409410715103
```

In [115]:

```
model = train3([50], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.009470

Iterations: 197

Function evaluations: 328

Gradient evaluations: 324

best_cost=0.009

Repetition 1: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.005514

Iterations: 335

Function evaluations: 470

Gradient evaluations: 463

best_cost=0.006

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.007878

Iterations: 135

Function evaluations: 234

Gradient evaluations: 227

best_cost=0.006

Repetition 3: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.081621

Iterations: 92

Function evaluations: 195

Gradient evaluations: 190

best_cost=0.006

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.081450

Iterations: 192

Function evaluations: 298

Gradient evaluations: 292

best_cost=0.006

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.006917

Iterations: 143

Function evaluations: 235

Gradient evaluations: 228

best_cost=0.006

Repetition 6: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.008208

Iterations: 201

Function evaluations: 328

Gradient evaluations: 324

best_cost=0.006

Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.006437

Iterations: 371

Function evaluations: 529

Gradient evaluations: 524

best_cost=0.006

Repetition 8: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.081728

Iterations: 126

Function evaluations: 218

Gradient evaluations: 210

best_cost=0.006

Repetition 9: Warning: Desired error not necessarily achieved due to precision loss.


Current function value: 0.081937

Iterations: 51

Function evaluations: 151

Gradient evaluations: 144

best_cost=0.006



Training error: 0.005512818694114685

Test error : 0.011191939003765583

In [116]:

```
model = train3([100], nreps=10, transfer=lambda: nn.ReLU())
nextplot()
plot1(X3, y3, label="train")
plot1(X3test, y3test, label="test")
plot1fit(torch.linspace(0, 13, 500).unsqueeze(1), model)
print("Training error:", F.mse_loss(y3, model(X3)).item())
print("Test error      :", F.mse_loss(y3test, model(X3test)).item())
```

Repetition 0: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.005350
Iterations: 297
Function evaluations: 415
Gradient evaluations: 408

best_cost=0.005

Repetition 1: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.081523
Iterations: 122
Function evaluations: 189
Gradient evaluations: 182

best_cost=0.005

Repetition 2: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.081756
Iterations: 78
Function evaluations: 157
Gradient evaluations: 153

best_cost=0.005

Repetition 3: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.005301
Iterations: 372
Function evaluations: 496
Gradient evaluations: 489

best_cost=0.005

Repetition 4: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.005127
Iterations: 382
Function evaluations: 470
Gradient evaluations: 462

best_cost=0.005

Repetition 5: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.005984
Iterations: 339
Function evaluations: 459
Gradient evaluations: 450

best_cost=0.005

Repetition 6: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.004477
Iterations: 542
Function evaluations: 703
Gradient evaluations: 696

best_cost=0.004

Repetition 7: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.080794

Iterations: 261

Function evaluations: 365

Gradient evaluations: 361

best_cost=0.004

Repetition 8: Warning: Desired error not necessarily achieved due to precision loss.

Current function value: 0.006943

Iterations: 196

Function evaluations: 323

Gradient evaluations: 317

best_cost=0.004

Repetition 9: Warning: Desired error not necessarily achieved due to precision loss.


Current function value: 0.006016

Iterations: 225

Function evaluations: 322

Gradient evaluations: 318

best_cost=0.004



Training error: 0.004473547451198101

Test error : 0.012758485972881317

