

a02-notebook

May 8, 2022

```
[ ]: # Student: Timur Carstensen  
# Student ID: 1722194
```

```
[ ]: #  
# IE 678 Deep Learning, University of Mannheim  
# Author: Rainer Gemulla
```

```
[ ]: import torch  
import torch.nn as nn  
import torchvision  
import matplotlib  
import matplotlib.pyplot as plt  
import os  
import numpy as np  
from sklearn.metrics import confusion_matrix  
  
from IPython import get_ipython  
from helper import *  
from util import nextplot  
  
%matplotlib inline
```

```
[ ]: # Use GPU if CUDA is available  
DATA_PATH = "data/"  
MODEL_PATH = "data/"  
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
```

```
[ ]: # To prevent from retraining models, you can save them to disk and load them  
# later on  
def save(model, filename):  
    torch.save(model.state_dict(), os.path.join(MODEL_PATH, filename))  
  
def load(model, filename):  
    model.load_state_dict(torch.load(os.path.join(MODEL_PATH, filename)))  
    return model
```

1 1 Convolutional Neural Networks

1.1 Load the data


```
[ ]: X, y, Xtest, ytest = load_dataset("fashionmnist")
class_dict = {
    0: "t-shirt/top",
    1: "trouser",
    2: "pullover",
    3: "dress",
    4: "coat",
    5: "sandal",
    6: "shirt",
    7: "sneaker",
    8: "bag",
    9: "ankle boot",
}
print(f"{len(X)} training examples")
print(f"{len(Xtest)} test examples")
```

60000 training examples
10000 test examples

```
[ ]: def show_image(x):
    "Show one (or multiple) 28x28 MNIST images as a gray-scale image."
    plt.imshow(x.reshape(-1, 28).cpu(), cmap="gray", interpolation="none")

    # Plot first 5 training examples. Each example consists of a 1x28x28 tensor
    # with values
    # in [0,1] and a label
    nextplot()
    show_image(X[:5])
    for i in range(5):
        print(f"Label of example i={i}: {class_dict[y[i].item()]})")
```

Label of example i=0: ankle boot
Label of example i=1: t-shirt/top
Label of example i=2: t-shirt/top
Label of example i=3: dress
Label of example i=4: t-shirt/top



1.2 1a+b Implement a CNN model

```
[ ]: # Here is a PyTorch version of logistic regression.
```

```
class LogisticRegression(nn.Module):
    def __init__(self, num_features):
        super().__init__()
        self.linear = nn.Linear(num_features, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        out = self.linear(x.float())
        out = self.sigmoid(out)
        return out
```

```
[ ]: # Implement a simple CNN
```

```
class SimpleCNN(nn.Module):
    def __init__(self):
        super().__init__()
        # Create the required layers/function and store them as instance_
        ↪ variables
        # YOUR CODE HERE

        self.layer1 = nn.Sequential(
            nn.Conv2d(kernel_size=3, in_channels=1, out_channels=32, stride=1, ↪
            ↪ padding=1, padding_mode="zeros"),
```

```

        nn.Sigmoid(),
        nn.MaxPool2d(kernel_size=2, stride=2),
    )
    self.fc1 = nn.Linear(in_features=6272, out_features=10)
    self.log_softmax = nn.LogSoftmax()

def forward(self, x):
    """
    Perform the forward pass.

    Parameters
    -----
    x: tensor of shape (batch_size, 1, 28, 28)

    Returns
    -----
    model output as a tensor of shape (batch_size, 10)
    """
    out = None
    # Use the layers/functions created above to compute model output
    # YOUR CODE HERE
    out = self.layer1(x)
    out = torch.flatten(out, 1)
    out = self.fc1(out)
    out = self.log_softmax(out)

    return out

```

```

[ ]: # here is a description of what you created (this uses the member variables)
print(SimpleCNN())

# SimpleCNN(
#   (layer1): Sequential(
#     (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
#     (1): Sigmoid()
#     (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ↵
#         ↪ceil_mode=False)
#   )
#   (fc1): Linear(in_features=6272, out_features=10, bias=True)
#   (log_softmax): LogSoftmax()
# )

```

```

SimpleCNN(
  (layer1): Sequential(
    (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): Sigmoid()

```

```

        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (fc1): Linear(in_features=6272, out_features=10, bias=True)
    (log_softmax): LogSoftmax(dim=None)
)

```

```

[ ]: # One way to see the parameters of your model is to look at its "state". Check
      ↳ that the
      # shapes of the parameters that you see here match your computations of task
      ↳ 1a).
model = SimpleCNN().to(DEVICE)
model.state_dict()

```

```

[ ]: OrderedDict([('layer1.0.weight',
      tensor([[[[-2.4413e-01,  4.2617e-02,  3.1756e-01],
                [-7.0482e-02,  1.7290e-02, -1.0355e-01],
                [ 3.1558e-02, -1.2980e-01,  3.2857e-01]]],
                [[[ 6.5351e-03,  9.9397e-02, -2.2767e-01],
                [ 2.9914e-01,  1.1743e-01,  4.4040e-02],
                [-2.8349e-01, -6.3990e-02,  2.0376e-01]]],
                [[[-1.0722e-01,  2.3458e-02,  5.1296e-02],
                [-6.3506e-03,  1.2136e-01, -1.3855e-01],
                [ 2.7447e-01,  4.4186e-02, -3.2704e-02]]],
                [[[-1.3704e-01, -2.0820e-01,  3.0723e-01],
                [-2.5863e-01,  2.2444e-01, -1.2951e-01],
                [ 1.4762e-01, -5.5286e-02,  1.5073e-01]]],
                [[[-3.0081e-01, -1.6919e-01,  9.5023e-02],
                [ 6.2922e-03, -2.3756e-01, -2.3372e-01],
                [-2.3056e-01,  1.1161e-02, -5.2626e-02]]],
                [[[-2.3484e-01,  1.2478e-02,  3.3158e-01],
                [ 1.5155e-01,  2.0131e-01, -3.0494e-01],
                [-1.5692e-01, -2.5777e-01, -7.3264e-02]]],
                [[[-1.1912e-01,  1.9668e-01, -2.6741e-01],
                [-2.2863e-02, -1.7096e-01,  6.7788e-02],

```

```

[ 6.1866e-02, -1.2586e-01, -1.2740e-01]]],

[[[-2.7602e-02, -1.3564e-01, -2.1127e-01],
 [ 2.1656e-01,  1.4129e-01, -2.4218e-01],
 [-1.9584e-01, -2.3433e-01, -1.7348e-01]]],

[[[-2.3539e-01, -1.3138e-01, -3.2553e-01],
 [-1.1223e-01, -2.0720e-01,  2.9110e-01],
 [ 3.9437e-02,  8.4036e-03, -4.0171e-02]]],

[[[-8.4452e-02,  2.9662e-01,  2.5323e-01],
 [ 2.7753e-01,  2.4395e-01, -1.2840e-01],
 [ 2.9421e-01, -1.8414e-01, -3.1143e-01]]],

[[[-2.2773e-01, -2.5044e-01, -1.3125e-01],
 [-7.8448e-02,  2.9281e-01,  1.4514e-01],
 [ 7.3863e-02, -2.6794e-01,  1.6106e-03]]],

[[[-2.8990e-01,  1.1630e-01,  5.3863e-02],
 [-3.3235e-02, -9.5065e-03, -3.1731e-01],
 [-8.7020e-02,  1.1119e-01, -1.8182e-02]]],

[[[-3.2994e-01, -1.9481e-01,  3.1867e-01],
 [-3.4428e-02, -1.4467e-01,  6.3393e-04],
 [ 5.3119e-02,  2.4644e-01, -2.4222e-01]]],

[[[-2.7880e-01, -1.1047e-01,  1.6090e-01],
 [-2.5582e-01, -9.9603e-02, -3.1559e-01],
 [-6.4307e-02,  2.6226e-01, -3.1198e-01]]],

[[[ 6.3477e-02, -2.0518e-01,  1.4477e-01],
 [-6.0590e-02,  2.6558e-01,  2.7426e-01],
 [-2.9511e-01,  3.3769e-02, -3.1349e-01]]],

[[[ 2.0599e-02,  2.1773e-01,  2.1587e-01],
 [ 3.1572e-01, -2.2507e-01,  6.9792e-02],
 [ 1.0981e-01, -1.7750e-01,  2.5679e-01]]],

```

```
[[[-2.2161e-01, 2.0771e-01, 1.7525e-01],  
 [-3.5824e-02, -1.1653e-01, -2.8805e-01],  
 [-9.0304e-02, -2.3229e-01, 1.6389e-01]]],
```

```
[[[ 2.8958e-01, -1.5212e-01, -2.3699e-01],  
 [ 2.6953e-01, 1.2830e-01, -3.9619e-02],  
 [ 3.1547e-01, 2.9752e-01, -1.5085e-01]]],
```

```
[[[ 1.0693e-01, 7.6142e-02, -1.3993e-01],  
 [ 1.6542e-01, -1.8272e-01, 3.0806e-01],  
 [ 1.4720e-01, -2.1153e-01, -2.9563e-03]]],
```

```
[[[-1.7642e-01, 9.0950e-02, 1.2180e-01],  
 [ 2.8606e-01, -3.1954e-01, -2.7676e-01],  
 [ 1.3660e-01, 1.0776e-01, -1.5727e-03]]],
```

```
[[[ 1.6720e-01, -2.6336e-04, -2.5182e-01],  
 [ 1.7061e-01, 3.0210e-01, 8.9676e-02],  
 [ 1.7778e-01, 3.0081e-01, -2.5432e-01]]],
```

```
[[[-1.4334e-01, 1.7685e-01, 5.0942e-02],  
 [-7.5547e-03, -1.5844e-01, -2.2489e-01],  
 [ 1.6404e-01, 1.5218e-01, -2.7796e-01]]],
```

```
[[[ 3.0824e-01, -2.1534e-02, 1.5150e-01],  
 [-3.0472e-01, -1.1949e-02, -2.8825e-01],  
 [-2.6855e-01, -1.9724e-01, -1.4083e-01]]],
```

```
[[[ 1.3439e-01, -3.7127e-02, 1.3735e-01],  
 [ 5.0808e-02, -2.3116e-01, -3.1062e-01],  
 [-9.1478e-02, -1.1906e-01, 2.5787e-01]]],
```

```
[[[-1.2887e-01, 1.8439e-01, -3.1071e-01],  
 [-3.3853e-03, 8.4458e-02, 9.3417e-03],  
 [-1.5631e-01, -2.5000e-01, -6.6180e-03]]],
```

```
[[[ 2.3939e-01, 2.9764e-02, 1.7813e-01],
```

```

        [-8.3289e-02,  2.5261e-01,  2.3511e-01],
        [ 1.1759e-01,  2.8765e-01,  1.3256e-01]]],

        [[[ 2.3241e-02,  7.4002e-02,  9.9316e-02],
         [ 1.6059e-01,  2.3090e-02,  1.1511e-01],
         [-2.6682e-01,  2.3910e-01, -2.2539e-01]]],

        [[[ 1.0598e-01,  2.5343e-01, -1.8034e-01],
         [-2.1721e-01,  2.0421e-01, -3.9024e-02],
         [ 1.2270e-01, -3.7205e-04,  1.2685e-02]]],

        [[[ 7.0892e-02,  2.6390e-01, -2.6142e-01],
         [ 1.5967e-01, -1.8723e-02, -3.2449e-01],
         [ 1.6712e-01, -1.9518e-01, -5.5869e-02]]],

        [[[ 1.2532e-03, -1.7499e-01, -1.0941e-01],
         [ 2.9381e-01, -1.0250e-01,  2.7956e-01],
         [-5.1640e-02,  2.9207e-02,  6.1875e-02]]],

        [[[ 2.4452e-01,  3.0876e-01, -4.8092e-03],
         [-1.9571e-01, -2.7417e-01, -1.7358e-01],
         [ 9.1399e-02,  1.1284e-01, -7.6799e-03]]],

        [[[ 2.2495e-01,  3.0138e-01, -7.7089e-02],
         [ 3.4538e-03,  1.6920e-01, -8.8056e-02],
         [ 2.3367e-01,  4.6525e-02,  2.5716e-01]]]])),
    ('layer1.0.bias',
     tensor([ 0.1437, -0.0331,  0.0052, -0.2324,  0.2480, -0.0151,
-0.0358,  0.2478,
-0.1060, -0.1052,  0.3299, -0.2489,  0.1990, -0.2528,
0.0044,  0.2342,
0.0279, -0.1444, -0.0194, -0.0461, -0.2956,  0.1159,
-0.0909, -0.0087,
-0.0306, -0.1698, -0.2644,  0.0666, -0.1278,  0.1081,
-0.0649,  0.3312])),
    ('fc1.weight',
     tensor([[ 0.0030, -0.0038, -0.0019, ..., -0.0062,  0.0048,
-0.0011],
[ 0.0067,  0.0064, -0.0086, ..., -0.0016,  0.0011,
0.0124],
[-0.0015,  0.0104, -0.0052, ..., -0.0099,  0.0007,

```



```
-0.0018],
        ...,
        [-0.0110,  0.0022, -0.0017, ..., -0.0086, -0.0051,
-0.0027],
        [ 0.0094, -0.0006,  0.0122, ...,  0.0076,  0.0049,
0.0052],
        [-0.0088, -0.0055, -0.0002, ..., -0.0109,  0.0005,
-0.0030]])),
        ('fc1.bias',
         tensor([ 0.0035, -0.0126, -0.0096,  0.0048,  0.0121,  0.0049,
-0.0088,  0.0044,
                 -0.0051,  0.0046]))))
```

```
[ ]: # Run the forward pass on the first 5 examples.
with torch.no_grad(): # tell torch to not compute backward graph
    print(model(X[:5]))
```

```
tensor([[ -2.0608, -2.0418, -2.5162, -2.9079, -2.1383, -2.2506, -1.8348, -2.6992,
          -2.2903, -2.8794],
        [-2.2168, -2.3191, -2.3916, -2.8283, -2.5097, -2.3808, -1.7170, -2.3541,
          -2.0668, -2.7011],
        [-2.1600, -2.2941, -2.4255, -2.9933, -2.4440, -2.3244, -1.7916, -2.4609,
          -2.0707, -2.4955],
        [-2.0819, -2.3726, -2.3657, -2.8566, -2.3902, -2.4340, -1.7798, -2.3896,
          -2.2686, -2.4349],
        [-2.2163, -2.3349, -2.3170, -3.0505, -2.3059, -2.3469, -1.7984, -2.3935,
          -2.0606, -2.6784]])
```

```
/var/folders/hq/pyn2nf8x3f94nbk8r749084r0000gn/T/ipykernel_72444/260300387.py:35
: UserWarning: Implicit dimension choice for log_softmax has been deprecated.
Change the call to include dim=X as an argument.
    out = self.log_softmax(out)
```

```
[ ]: # Plot training (or test) examples, their correct labels, and the most likely
      ↪ model
      # predictions. Do not be worried if your (untrained) model always seems to
      ↪ predict the
      # same class.
@torch.no_grad()
def mnist_predict(model, start, end=None, use_test_data=False):
    if end is None:
        end = start + 1
    if use_test_data:
        images = Xtest[start:end].to(DEVICE)
        labels = ytest[start:end].to(DEVICE)
    else:
        images = X[start:end].to(DEVICE)
```

```

        labels = y[start:end].to(DEVICE)
    nextplot()
    show_image(images)
    print("    Labels:", [class_dict[label.item()] for label in labels])
    out = model(images)
    _, yhat = torch.max(out, 1)
    print("Predictions:", [class_dict[pred.item()] for pred in yhat])

# first 5 examples from training + predictions
mnist_predict(model, 0, 5)

```

Labels: ['ankle boot', 't-shirt/top', 't-shirt/top', 'dress',
't-shirt/top']


Predictions: ['shirt', 'shirt', 'shirt', 'shirt', 'shirt']

/var/folders/hq/pyn2nf8x3f94nbk8r749084r0000gn/T/ipykernel_72444/260300387.py:35

: UserWarning: Implicit dimension choice for log_softmax has been deprecated.

Change the call to include dim=X as an argument.

```
out = self.log_softmax(out)
```



1.3 1c Evaluate model performance

```
[ ]: # test model
@torch.no_grad()
def mnist_test(model, batch_size=100, reshape_id=False):
    """
    Function to test your CNN on test data

    Parameters
    -----

    model: trained CNN from task 1a
    batch_size: size of batch for dataloader
    reshape_id: Reshape images to a 1d vectors (allows use of models other than
    ↪ CNNs
                such as fully-connected FNNs)

    Returns
    -----
    accuracy of input model
    """
    correct = 0 # number of correct predictions
    total = 0 # total number of examples
    model.eval() # set layers like dropout and batch norm to eval mode

    # Create test data loader
    if reshape_id:
        dataset = torch.utils.data.TensorDataset(Xtest.reshape(len(Xtest), -1),
        ↪ ytest)
    else:
        dataset = torch.utils.data.TensorDataset(Xtest, ytest)
    test_loader = torch.utils.data.DataLoader(
        dataset, batch_size=batch_size, shuffle=False
    )
    # confusion_matrix = np.zeros(shape=(10,10))
    pred = list()
    y = list()
    # Loop over data
    for batch in test_loader:
        # YOUR CODE HERE
        # Update correct and total using the examples in the batch. To
        ↪ understand what a
            # DataLoader does, have a look at the contents of "batch" before you
        ↪ start.
        images, labels = batch
        # calculate outputs by running images through the network
        outputs = model(images)
```

```

        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        pred.append(predicted)
        y.append(labels)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        # confusion matrix over classes

M = confusion_matrix(torch.cat(y), torch.cat(pred))
accuracy = (correct / total) * 100
print(f"Accuracy on {total} test images: {accuracy:.2f} %")
return accuracy, M

```

```

[ ]: # Test your code. Output should be: 6.97%
class DummyModel(nn.Module):
    def forward(self, x):
        return x.reshape(len(x), -1)[: , 200:210] * torch.arange(10).to(DEVICE)

mnist_test(DummyModel().to(DEVICE))[0]

```

Accuracy on 10000 test images: 6.97 %

```
[ ]: 6.97
```

1.4 1d Train a model

```

[ ]: # train model
def mnist_train(
    model, num_epochs=5, learning_rate=0.001, batch_size=100, reshape_1d=False
):
    """
    Function to train the provided CNN network.

    Parameters
    -----
    model: the model to train
    num_epochs: number of epochs to train
    learning_rate: learning rate to use
    batch_size: size of batch for data loader
    reshape_1d: Reshape images to a 1d vectors (allows use of models other than
    ↪ CNNs
    such as fully-connected FNNs)
    """
    # YOUR CODE HERE

    if reshape_1d:

```

```

        dataset = torch.utils.data.TensorDataset(X.reshape(len(X), -1), y)
    else:
        dataset = torch.utils.data.TensorDataset(X, y)
    test_loader = torch.utils.data.DataLoader(
        dataset, batch_size=batch_size, shuffle=False
    )

    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
    criterion = nn.NLLLoss()
    i = 0

    for epoch in range(num_epochs): # loop over the dataset multiple times

        running_loss = 0.0

        for batch in test_loader:
            # get the inputs; data is a list of [inputs, labels]
            inputs, labels = batch

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            i += 1
            if i % 2000 == 1999: # print every 2000 mini-batches
                print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.
↪3f}')

                running_loss = 0.0

    return model

```

1.5 1e Train and evaluate the simple CNN model

```

[ ]: # load model
load(model, "simple_cnn.pt")

```

```

[ ]: SimpleCNN(
    (layer1): Sequential(

```

```

        (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): Sigmoid()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (fc1): Linear(in_features=6272, out_features=10, bias=True)
    (log_softmax): LogSoftmax(dim=None)
)

```

```

[ ]: # Train a model.
model = mnist_train(SimpleCNN().to(DEVICE))

```

```

/var/folders/hq/pyn2nf8x3f94nbk8r749084r0000gn/T/ipykernel_72444/260300387.py:35
: UserWarning: Implicit dimension choice for log_softmax has been deprecated.
Change the call to include dim=X as an argument.
    out = self.log_softmax(out)

[4, 2000] loss: 0.027

```

```

[ ]: # persisting the model
save(model, "simple_cnn.pt")

```

```

[ ]: # Test a model. The simple CNN should perform much better after training.
# model = load(SimpleCNN(), "simple_cnn.pt").to(DEVICE)
_, M = mnist_test(model)
plt.imshow(M, origin="upper")
for ij, v in np.ndenumerate(M):
    i, j = ij
    plt.text(j, i, str(v), color="white", ha="center", va="center")
plt.xlabel("predicted")
plt.ylabel("true")
plt.colorbar()
plt.savefig("tex/assets/conf_mat.png", dpi=1000)
acc = M.diagonal()/M.sum(axis=1)
from util import plot_matrix

out = {v:acc[i] for i, (k, v) in enumerate(class_dict.items())}
plot_matrix(M, colnames=class_dict.items())

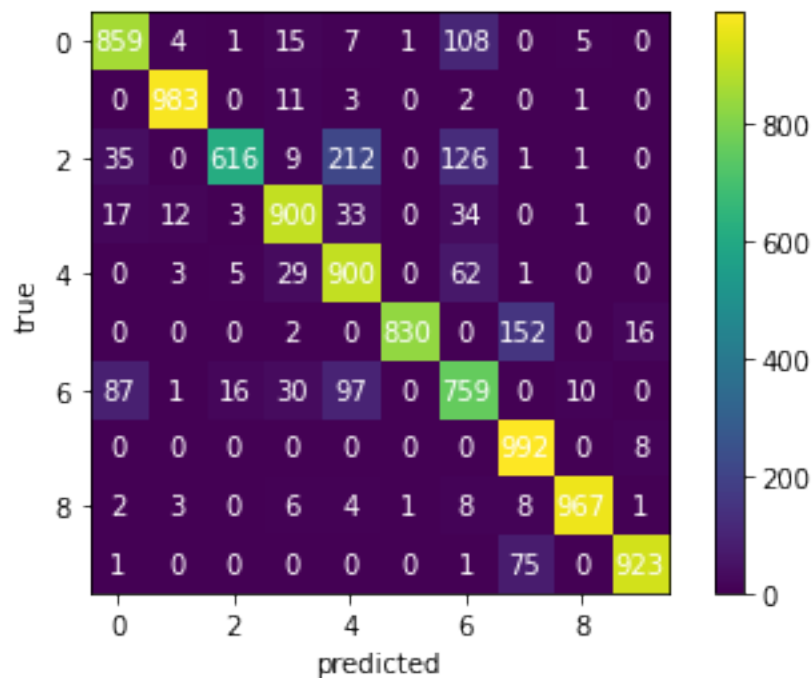
```

```

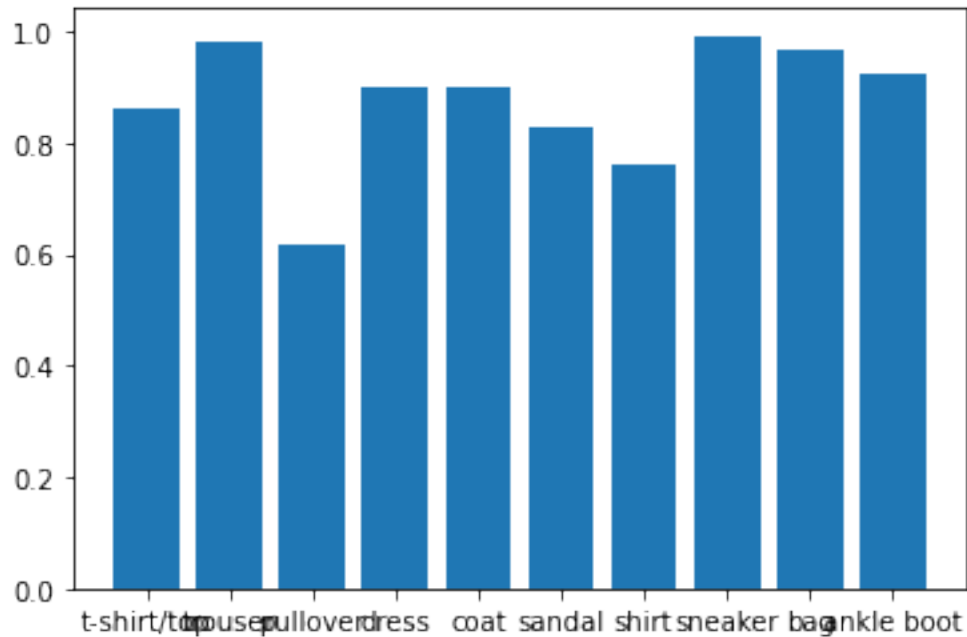
/var/folders/hq/pyn2nf8x3f94nbk8r749084r0000gn/T/ipykernel_72444/260300387.py:35
: UserWarning: Implicit dimension choice for log_softmax has been deprecated.
Change the call to include dim=X as an argument.
    out = self.log_softmax(out)

Accuracy on 10000 test images: 87.29 %

```



```
[ ]: nextplot()
plt.bar(range(len(out)), list(out.values()), align='center')
plt.xticks(range(len(out)), list(out.keys()))
plt.show()
plt.savefig("tex/assets/accuracies_per_class.png", dpi=200)
```



<Figure size 432x288 with 0 Axes>

1.6 1f Sandbox

```
[ ]: # try adding more layers, regularization, etc.
class MyCnn(nn.Module):
    def __init__(self):
        super(MyCnn, self).__init__()
        # YOUR CODE HERE

    def forward(self, x):
        out = None
        # YOUR CODE HERE

        return out
```

```
[ ]: my_cnn = MyCnn().to(DEVICE)
```



```
[ ]: # train
mnist_train(my_cnn)
# save(my_cnn, "my_cnn.pt")
```

```
[ ]: # test
# my_cnn = load(my_cnn, "my_cnn.pt")
mnist_test(my_cnn)
```

2 2 Recurrent Neural Networks and Pretraining

2.1 Load and preprocess the data

```
[ ]: # Load review and label data
with open(os.path.join(DATA_PATH, "reviews_small.txt")) as f:
    reviews_lines = f.readlines()

with open(os.path.join(DATA_PATH, "labels_small.txt")) as f:
    label_lines = f.readlines()

print(reviews_lines[0])
print(label_lines[0])
```

bromwell high is a cartoon comedy . it ran at the same time as some other programs about school life such as teachers . my years in the teaching profession lead me to believe that bromwell high s satire is much closer to reality than is teachers . the scramble to survive financially the insightful students who can see right through their pathetic teachers pomp the pettiness of the whole situation all remind me of the schools i knew and their students . when i saw the episode in which a student repeatedly tried to burn down the school i immediately recalled at high . a classic line inspector i m here to sack one of your teachers . student welcome to bromwell high . i expect that many adults of my age think that bromwell high is far fetched . what a pity that it isn t

positive

```
[ ]: # Remove punctuations from the reviews and split words
raw_reviews, words, labels = reviews_preprocess(reviews_lines, label_lines)
print(raw_reviews[0])
print("First ten words: ", words[:10])
print("First label: ", labels[0])
```

bromwell high is a cartoon comedy it ran at the same time as some other programs about school life such as teachers my years in the teaching profession lead me to believe that bromwell high s satire is much closer to reality than is teachers the scramble to survive financially the insightful students who can see

right through their pathetic teachers pomp the pettiness of the whole situation
all remind me of the schools i knew and their students when i saw the episode in
which a student repeatedly tried to burn down the school i immediately recalled
at high a classic line inspector i m here to sack one of your teachers student
welcome to bromwell high i expect that many adults of my age think that bromwell
high is far fetched what a pity that it isn t
First ten words: ['bromwell', 'high', 'is', 'a', 'cartoon', 'comedy', 'it',
'ran', 'at', 'the']
First label: 1

```
[ ]: # Determine an integer id for each unique word
word_ids = reviews_create_word_ids(words)
print(word_ids.get("the"))
print(word_ids.get("movie"))
```

1
16

```
[ ]: # Encode each word in the review by its unique identifier
encoded_reviews = reviews_encode(word_ids, raw_reviews)
print(raw_reviews[0])
print(encoded_reviews[0])
```

bromwell high is a cartoon comedy it ran at the same time as some other programs
about school life such as teachers my years in the teaching profession lead me
to believe that bromwell high s satire is much closer to reality than is
teachers the scramble to survive financially the insightful students who can see
right through their pathetic teachers pomp the pettiness of the whole situation
all remind me of the schools i knew and their students when i saw the episode in
which a student repeatedly tried to burn down the school i immediately recalled
at high a classic line inspector i m here to sack one of your teachers student
welcome to bromwell high i expect that many adults of my age think that bromwell
high is far fetched what a pity that it isn t

[10455, 307, 6, 3, 1177, 202, 8, 2217, 33, 1, 168, 56, 15, 49, 85, 8269, 43,
422, 122, 140, 15, 3151, 59, 144, 9, 1, 5230, 5946, 452, 72, 5, 260, 12, 10455,
307, 13, 2017, 6, 73, 2765, 5, 689, 76, 6, 3151, 1, 19565, 5, 1706, 6897, 1,
5947, 1707, 36, 52, 68, 211, 143, 63, 1390, 3151, 14816, 1, 19566, 4, 1, 221,
755, 31, 2710, 72, 4, 1, 5948, 10, 729, 2, 63, 1707, 54, 10, 208, 1, 321, 9, 64,
3, 1578, 3922, 737, 5, 2843, 187, 1, 422, 10, 1246, 9217, 33, 307, 3, 380, 322,
5949, 10, 135, 136, 5, 9218, 30, 4, 134, 3151, 1578, 2480, 5, 10455, 307, 10,
528, 12, 113, 1839, 4, 59, 676, 103, 12, 10455, 307, 6, 227, 4097, 48, 3, 2169,
12, 8, 231, 21]

```
[ ]: # Padding/truncating all reviews to the same length. Although this isn't
↳ strictly
# necessary, it facilitates batch processing: all inputs of a batch need to
↳ have the
```

```
# same length.
sequence_length = 200
padded_reviews = reviews_pad(encoded_reviews, sequence_length)
print(padded_reviews[0])
```

```
[ 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0
10455 307 6 3 1177 202 8 2217 33 1 168 56
 15 49 85 8269 43 422 122 140 15 3151 59 144
 9 1 5230 5946 452 72 5 260 12 10455 307 13
2017 6 73 2765 5 689 76 6 3151 1 19565 5
1706 6897 1 5947 1707 36 52 68 211 143 63 1390
3151 14816 1 19566 4 1 221 755 31 2710 72 4
 1 5948 10 729 2 63 1707 54 10 208 1 321
 9 64 3 1578 3922 737 5 2843 187 1 422 10
1246 9217 33 307 3 380 322 5949 10 135 136 5
9218 30 4 134 3151 1578 2480 5 10455 307 10 528
 12 113 1839 4 59 676 103 12 10455 307 6 227
4097 48 3 2169 12 8 231 21]
```

```
[ ]: # Split dataset into 80% training, 10% test, and 10% validation Dataset
train_x, train_y, valid_x, valid_y, test_x, test_y = reviews_split(
    padded_reviews, labels
)
print(len(train_y), len(valid_y), len(test_y))
```

```
3200 400 400
```

```
[ ]: # Create data loaders for training
train_loader, valid_loader, test_loader = reviews_create_dataloaders(
    train_x, train_y, valid_x, valid_y, test_x, test_y
)
```

```
[ ]: # Here is an example how to use the train and test functions. Note that logistic
# regression is a bogus model when used like this (since its assigns weights to
# positions, but not word ids). So results will be bad.
model = LogisticRegression(sequence_length).to(DEVICE)
reviews_train(model, train_loader, valid_loader, epochs=3, device=DEVICE)
reviews_test(model, test_loader, device=DEVICE)
```

```
Starting epoch 1
```

```
Epoch: 1/ 3    Batch: 5    Batch loss: 55.862564    Val loss: 52.007381 Val
acc: 0.470000
```

```
Epoch: 1/ 3    Batch: 10    Batch loss: 44.230236    Val loss: 50.456723 Val
acc: 0.490000
```

Epoch: 1/ 3	Batch: 15	Batch loss: 50.000000	Val loss: 50.462179	Val
acc: 0.490000				
Epoch: 1/ 3	Batch: 20	Batch loss: 44.000000	Val loss: 50.634626	Val
acc: 0.490000				
Epoch: 1/ 3	Batch: 25	Batch loss: 32.900143	Val loss: 52.853641	Val
acc: 0.470000				
Epoch: 1/ 3	Batch: 30	Batch loss: 52.671864	Val loss: 51.840658	Val
acc: 0.480000				
Epoch: 1/ 3	Batch: 35	Batch loss: 62.000000	Val loss: 52.332016	Val
acc: 0.470000				
Epoch: 1/ 3	Batch: 40	Batch loss: 53.558922	Val loss: 51.445327	Val
acc: 0.480000				
Epoch: 1/ 3	Batch: 45	Batch loss: 48.000038	Val loss: 51.347229	Val
acc: 0.480000				
Epoch: 1/ 3	Batch: 50	Batch loss: 58.108727	Val loss: 51.581363	Val
acc: 0.480000				
Epoch: 1/ 3	Batch: 55	Batch loss: 54.000000	Val loss: 51.581363	Val
acc: 0.480000				
Epoch: 1/ 3	Batch: 60	Batch loss: 46.955475	Val loss: 51.581358	Val
acc: 0.480000				

96

Finished epoch 1. Average batch loss: 49.876608073711395. Average validation loss: 51.51032213370005

Starting epoch 2

Epoch: 2/ 3	Batch: 65	Batch loss: 54.000000	Val loss: 50.764032	Val
acc: 0.490000				
Epoch: 2/ 3	Batch: 70	Batch loss: 46.027721	Val loss: 51.250000	Val
acc: 0.487500				
Epoch: 2/ 3	Batch: 75	Batch loss: 52.763466	Val loss: 50.798164	Val
acc: 0.487500				
Epoch: 2/ 3	Batch: 80	Batch loss: 46.976650	Val loss: 51.250000	Val
acc: 0.487500				
Epoch: 2/ 3	Batch: 85	Batch loss: 52.000000	Val loss: 51.250000	Val
acc: 0.487500				
Epoch: 2/ 3	Batch: 90	Batch loss: 44.125820	Val loss: 51.478319	Val
acc: 0.477500				
Epoch: 2/ 3	Batch: 95	Batch loss: 42.000000	Val loss: 51.478445	Val
acc: 0.477500				
Epoch: 2/ 3	Batch: 100	Batch loss: 56.000000	Val loss: 51.692595	Val
acc: 0.480000				
Epoch: 2/ 3	Batch: 105	Batch loss: 38.283012	Val loss: 52.004820	Val
acc: 0.477500				
Epoch: 2/ 3	Batch: 110	Batch loss: 42.000000	Val loss: 51.599065	Val
acc: 0.482500				
Epoch: 2/ 3	Batch: 115	Batch loss: 45.428814	Val loss: 51.441724	Val
acc: 0.482500				
Epoch: 2/ 3	Batch: 120	Batch loss: 52.000000	Val loss: 51.477280	Val
acc: 0.482500				

Epoch: 2/ 3 Batch: 125 Batch loss: 40.000000 Val loss: 51.572995 Val
acc: 0.482500
104

Finished epoch 2. Average batch loss: 49.45479238033295. Average validation
loss: 51.38903372104351

Starting epoch 3

Epoch: 3/ 3 Batch: 130 Batch loss: 50.000000 Val loss: 51.572995 Val
acc: 0.482500

Epoch: 3/ 3 Batch: 135 Batch loss: 57.750259 Val loss: 50.593139 Val
acc: 0.487500

Epoch: 3/ 3 Batch: 140 Batch loss: 50.000000 Val loss: 51.007401 Val
acc: 0.487500

Epoch: 3/ 3 Batch: 145 Batch loss: 60.000000 Val loss: 50.301833 Val
acc: 0.492500

Epoch: 3/ 3 Batch: 150 Batch loss: 51.195168 Val loss: 50.301455 Val
acc: 0.492500

Epoch: 3/ 3 Batch: 155 Batch loss: 52.000000 Val loss: 50.271038 Val
acc: 0.492500

Epoch: 3/ 3 Batch: 160 Batch loss: 54.000942 Val loss: 50.500000 Val
acc: 0.495000

Epoch: 3/ 3 Batch: 165 Batch loss: 51.285694 Val loss: 50.500000 Val
acc: 0.495000

Epoch: 3/ 3 Batch: 170 Batch loss: 60.000000 Val loss: 50.500000 Val
acc: 0.495000

Epoch: 3/ 3 Batch: 175 Batch loss: 39.232018 Val loss: 50.030633 Val
acc: 0.497500

Epoch: 3/ 3 Batch: 180 Batch loss: 58.000000 Val loss: 50.030633 Val
acc: 0.497500

Epoch: 3/ 3 Batch: 185 Batch loss: 58.000000 Val loss: 50.431529 Val
acc: 0.495000

Epoch: 3/ 3 Batch: 190 Batch loss: 60.000000 Val loss: 50.116792 Val
acc: 0.495000

104

Finished epoch 3. Average batch loss: 49.30076479911804. Average validation
loss: 50.473649795238785

Test loss: 49.750

Test accuracy: 0.502

2.2 2a Define your model

```
[ ]: # Create an LSTM for sentiment analysis
class SimpleLSTM(nn.Module):
    """
    The RNN model that will be used to perform sentiment analysis.
    """

    def __init__(
```

```

self,
vocab_size,
embedding_dim,
hidden_dim,
num_layers=1,
lstm_dropout_prob=0.5,
dropout_prob=0.3,
):
    """
    Initialize the model by setting up the layers

    Parameters
    -----
    vocab_size: number of unique words in the reviews
    embeddings_dim: size of the embeddings
    hidden_dim: dimension of the LSTM output
    num_layers: number of LSTM layers
    lstm_dropout_prob: dropout applied between the LSTM layers
    dropout_prob: dropout applied before the fully connected layer
    """
    super().__init__()

    self.num_layers = num_layers
    self.hidden_dim = hidden_dim

    self.embedding = nn.Embedding(
        num_embeddings=vocab_size,
        embedding_dim=embedding_dim
    )

    self.lstm = nn.LSTM(
        input_size=embedding_dim,
        hidden_size=hidden_dim,
        num_layers=num_layers,
        batch_first=True,
        dropout=lstm_dropout_prob
    )

    self.dropout = nn.Dropout(
        p=dropout_prob,
        inplace=False
    )

    self.fc = nn.Linear(
        in_features=hidden_dim,
        out_features=1
    )

```

```

self.sigmoid = nn.Sigmoid()

# YOUR CODE HERE

def forward(self, x):
    """
    Perform a forward pass of our model on some input and hidden state.

    Parameters
    -----
    x: batch as a (batch_size, sequence_length) tensor

    Returns
    -----
    Probability of positive class.
    """
    # init hidden layer, which is needed for the LSTM
    batch_size = len(x)
    hidden = self.init_hidden(batch_size)

    # YOUR CODE HERE
    embedding = self.embedding(x)
    out, (hidden, cell_state) = self.lstm(embedding, hidden)
    out = self.dropout(out[:, -1, :])
    out = self.fc(out)
    out = self.sigmoid(out)
    return out

def init_hidden(self, batch_size):
    """
    Initialize hidden state.

    Returns
    -----
    Empty hidden LSTM state.
    """

    # Create two new tensors with sizes num_layers x batch_size x
    ↪ hidden_dim,
    # initialized to zero, for hidden state and cell state of LSTM
    weight = next(self.parameters()) # only used to determine device

    hidden = (
        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_(),

```

```

        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_(),
    )

    return hidden

```

```

[ ]: # Test model setup
lstm_model = SimpleLSTM(1, 10, 32, 2, 0, 0).to(DEVICE)
print(lstm_model)

# SimpleLSTM(
#   (embedding): Embedding(1, 10)
#   (lstm): LSTM(10, 32, num_layers=2, batch_first=True, dropout=0.0)
#   (dropout): Dropout(p=0.0, inplace=False)
#   (fc): Linear(in_features=32, out_features=1, bias=True)
#   (sigmoid): Sigmoid()
# )

```

```

SimpleLSTM(
  (embedding): Embedding(1, 10)
  (lstm): LSTM(10, 32, num_layers=2, batch_first=True)
  (dropout): Dropout(p=0, inplace=False)
  (fc): Linear(in_features=32, out_features=1, bias=True)
  (sigmoid): Sigmoid()
)

```

```

[ ]: # Test forward function.
# dummy data
dummy_data = torch.zeros(train_loader.batch_size, sequence_length).long().
    ↪to(DEVICE)
# fix model parameters
for key in lstm_model.state_dict():
    lstm_model.state_dict()[key][:] = 0.1
print(lstm_model(dummy_data).reshape(10, -1))
# Output after reshape should be the following tensor
#tensor([[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
#        [0.9643, 0.9643, 0.9643, 0.9643, 0.9643]], device='cuda or cpu',
#        grad_fn=<ViewBackward>)

```

```

tensor([[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],

```



```

[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643],
[0.9643, 0.9643, 0.9643, 0.9643, 0.9643]],
grad_fn=<ReshapeAliasBackward0>)

```

2.3 2b Train and evaluate the model

```

[ ]: # Instantiate the model w/ hyperparams
vocab_size = len(word_ids) + 1 # +1 for the 0 padding
embedding_dim = 100
hidden_dim = 64
num_layers = 1

# this may raise a warning when num_layers=1 (which is fine)
# make sure to reinizialize the model if you want to train multiple times
lstm_model = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers).
    ↪to(DEVICE)

```

```

[ ]: # Fit and evaluate a model (without pretrained embeddings)
n_epochs = 5
# YOUR CODE HERE
reviews_train(lstm_model, train_loader, valid_loader, epochs=5, device=DEVICE)

```

Starting epoch 1

Epoch: 1/ 5	Batch: 5	Batch loss: 0.679525	Val loss: 0.703965	Val acc: 0.495000
Epoch: 1/ 5	Batch: 10	Batch loss: 0.698851	Val loss: 0.695553	Val acc: 0.517500
Epoch: 1/ 5	Batch: 15	Batch loss: 0.705797	Val loss: 0.695675	Val acc: 0.512500
Epoch: 1/ 5	Batch: 20	Batch loss: 0.708614	Val loss: 0.694467	Val acc: 0.550000
Epoch: 1/ 5	Batch: 25	Batch loss: 0.667098	Val loss: 0.693591	Val acc: 0.547500
Epoch: 1/ 5	Batch: 30	Batch loss: 0.688021	Val loss: 0.694176	Val acc: 0.530000
Epoch: 1/ 5	Batch: 35	Batch loss: 0.679528	Val loss: 0.693740	Val acc: 0.540000
Epoch: 1/ 5	Batch: 40	Batch loss: 0.673149	Val loss: 0.693195	Val acc: 0.552500
Epoch: 1/ 5	Batch: 45	Batch loss: 0.683008	Val loss: 0.691528	Val acc: 0.547500

Epoch: 1/ 5 Batch: 50 Batch loss: 0.701330 Val loss: 0.689682 Val acc:
 0.550000
 Epoch: 1/ 5 Batch: 55 Batch loss: 0.668185 Val loss: 0.687032 Val acc:
 0.547500
 Epoch: 1/ 5 Batch: 60 Batch loss: 0.677919 Val loss: 0.685461 Val acc:
 0.565000
 96
 Finished epoch 1. Average batch loss: 0.6865371307358146. Average validation
 loss: 0.693172030771772
 Starting epoch 2
 Epoch: 2/ 5 Batch: 65 Batch loss: 0.672716 Val loss: 0.684752 Val acc:
 0.555000
 Epoch: 2/ 5 Batch: 70 Batch loss: 0.634400 Val loss: 0.682325 Val acc:
 0.562500
 Epoch: 2/ 5 Batch: 75 Batch loss: 0.627135 Val loss: 0.677084 Val acc:
 0.597500
 Epoch: 2/ 5 Batch: 80 Batch loss: 0.644966 Val loss: 0.679041 Val acc:
 0.555000
 Epoch: 2/ 5 Batch: 85 Batch loss: 0.568219 Val loss: 0.812599 Val acc:
 0.535000
 Epoch: 2/ 5 Batch: 90 Batch loss: 0.628705 Val loss: 0.682299 Val acc:
 0.557500
 Epoch: 2/ 5 Batch: 95 Batch loss: 0.622430 Val loss: 0.671059 Val acc:
 0.600000
 Epoch: 2/ 5 Batch: 100 Batch loss: 0.538470 Val loss: 0.666089 Val acc:
 0.600000
 Epoch: 2/ 5 Batch: 105 Batch loss: 0.653860 Val loss: 0.669573 Val acc:
 0.605000
 Epoch: 2/ 5 Batch: 110 Batch loss: 0.577353 Val loss: 0.665001 Val acc:
 0.612500
 Epoch: 2/ 5 Batch: 115 Batch loss: 0.587399 Val loss: 0.657142 Val acc:
 0.615000
 Epoch: 2/ 5 Batch: 120 Batch loss: 0.686283 Val loss: 0.650623 Val acc:
 0.625000
 Epoch: 2/ 5 Batch: 125 Batch loss: 0.596217 Val loss: 0.715745 Val acc:
 0.575000
 104
 Finished epoch 2. Average batch loss: 0.6149493930861354. Average validation
 loss: 0.6856409878684924
 Starting epoch 3
 Epoch: 3/ 5 Batch: 130 Batch loss: 0.498110 Val loss: 0.642683 Val acc:
 0.632500
 Epoch: 3/ 5 Batch: 135 Batch loss: 0.568852 Val loss: 0.667370 Val acc:
 0.610000
 Epoch: 3/ 5 Batch: 140 Batch loss: 0.532824 Val loss: 0.660348 Val acc:
 0.622500
 Epoch: 3/ 5 Batch: 145 Batch loss: 0.475732 Val loss: 0.659632 Val acc:
 0.637500

Epoch: 3/ 5	Batch: 150	Batch loss: 0.547082	Val loss: 0.698349	Val acc: 0.590000
Epoch: 3/ 5	Batch: 155	Batch loss: 0.509110	Val loss: 0.642167	Val acc: 0.652500
Epoch: 3/ 5	Batch: 160	Batch loss: 0.507828	Val loss: 0.656309	Val acc: 0.645000
Epoch: 3/ 5	Batch: 165	Batch loss: 0.442918	Val loss: 0.642163	Val acc: 0.647500
Epoch: 3/ 5	Batch: 170	Batch loss: 0.548849	Val loss: 0.642707	Val acc: 0.650000
Epoch: 3/ 5	Batch: 175	Batch loss: 0.440895	Val loss: 0.635169	Val acc: 0.635000
Epoch: 3/ 5	Batch: 180	Batch loss: 0.491875	Val loss: 0.649510	Val acc: 0.650000
Epoch: 3/ 5	Batch: 185	Batch loss: 0.967413	Val loss: 0.647366	Val acc: 0.637500
Epoch: 3/ 5	Batch: 190	Batch loss: 0.559213	Val loss: 0.735722	Val acc: 0.542500

104

Finished epoch 3. Average batch loss: 0.522889809217304. Average validation loss: 0.6599610298871994

Starting epoch 4

Epoch: 4/ 5	Batch: 195	Batch loss: 0.521300	Val loss: 0.649883	Val acc: 0.645000
Epoch: 4/ 5	Batch: 200	Batch loss: 0.353082	Val loss: 0.648531	Val acc: 0.652500
Epoch: 4/ 5	Batch: 205	Batch loss: 0.409361	Val loss: 0.635625	Val acc: 0.670000
Epoch: 4/ 5	Batch: 210	Batch loss: 0.418146	Val loss: 0.628939	Val acc: 0.662500
Epoch: 4/ 5	Batch: 215	Batch loss: 0.571482	Val loss: 0.676091	Val acc: 0.610000
Epoch: 4/ 5	Batch: 220	Batch loss: 0.525052	Val loss: 0.682291	Val acc: 0.632500
Epoch: 4/ 5	Batch: 225	Batch loss: 0.456031	Val loss: 0.654979	Val acc: 0.665000
Epoch: 4/ 5	Batch: 230	Batch loss: 0.340121	Val loss: 0.654314	Val acc: 0.652500
Epoch: 4/ 5	Batch: 235	Batch loss: 0.343296	Val loss: 0.661900	Val acc: 0.655000
Epoch: 4/ 5	Batch: 240	Batch loss: 0.322669	Val loss: 0.652139	Val acc: 0.665000
Epoch: 4/ 5	Batch: 245	Batch loss: 0.371951	Val loss: 0.722962	Val acc: 0.642500
Epoch: 4/ 5	Batch: 250	Batch loss: 0.364596	Val loss: 0.744807	Val acc: 0.665000
Epoch: 4/ 5	Batch: 255	Batch loss: 0.424592	Val loss: 0.642369	Val acc: 0.662500

104

Finished epoch 4. Average batch loss: 0.42475875513628125. Average validation loss: 0.6657560762877648

Starting epoch 5

Epoch: 5/ 5 Batch: 260 Batch loss: 0.520416 Val loss: 0.673290 Val acc: 0.657500

Epoch: 5/ 5 Batch: 265 Batch loss: 0.435737 Val loss: 0.651521 Val acc: 0.667500

Epoch: 5/ 5 Batch: 270 Batch loss: 0.303476 Val loss: 0.634039 Val acc: 0.672500

Epoch: 5/ 5 Batch: 275 Batch loss: 0.378886 Val loss: 0.698585 Val acc: 0.647500

Epoch: 5/ 5 Batch: 280 Batch loss: 0.423104 Val loss: 0.630883 Val acc: 0.687500

Epoch: 5/ 5 Batch: 285 Batch loss: 0.296592 Val loss: 0.645418 Val acc: 0.687500

Epoch: 5/ 5 Batch: 290 Batch loss: 0.341979 Val loss: 0.657843 Val acc: 0.682500

Epoch: 5/ 5 Batch: 295 Batch loss: 0.260322 Val loss: 0.674640 Val acc: 0.675000

Epoch: 5/ 5 Batch: 300 Batch loss: 0.216686 Val loss: 0.748435 Val acc: 0.657500

Epoch: 5/ 5 Batch: 305 Batch loss: 0.300631 Val loss: 0.652457 Val acc: 0.672500

Epoch: 5/ 5 Batch: 310 Batch loss: 0.452203 Val loss: 0.664092 Val acc: 0.647500

Epoch: 5/ 5 Batch: 315 Batch loss: 0.334228 Val loss: 0.669618 Val acc: 0.675000

Epoch: 5/ 5 Batch: 320 Batch loss: 0.356790 Val loss: 0.666637 Val acc: 0.690000

104

Finished epoch 5. Average batch loss: 0.3577388192061335. Average validation loss: 0.6667274890037683

```
[ ]: reviews_test(lstm_model, test_loader, device=DEVICE)
```

Test loss: 0.918

Test accuracy: 0.682

2.4 2c Load pretrained word embeddings

```
[ ]: @torch.no_grad()
def reviews_load_embeddings(
    embedding_layer, word_ids, pretrained_embeddings_file="data/word-embeddings.
    ↪txt"
):
    """Load pretrained embeddings into an embedding layer.
```

Updates the weights of the embedding layer with the embeddings given in the provided word embeddings file.

Parameters

embedding_layer: torch.nn.Embedding used in the model

word_ids: dictionary mapping each word to its unique identifier

pretrained_embeddings_file: path to the file containing pretrained embeddings

```
"""
print("Initializing embedding layer with pretrained word embeddings...")
embeddings_index = dict()
words_initialized = 0
with open(pretrained_embeddings_file, encoding="utf8") as f:
    for line in f:
        values = line.split()
        word = values[0]
        encoded_word = word_ids.get(word)
        if encoded_word is not None:
            words_initialized += 1
            embedding_layer.weight[encoded_word, :] = torch.from_numpy(
                np.asarray(values[1:], dtype="float32")
            )
print(
    "Initialized {}/{} word embeddings".format(
        words_initialized, embedding_layer.num_embeddings
    )
)
```

```
[ ]: # Try it
test_embeddings = nn.Embedding(len(word_ids) + 1, 100).to(DEVICE)
reviews_load_embeddings(test_embeddings, word_ids)
print(test_embeddings(torch.LongTensor([word_ids.get("movie")]).to(DEVICE)))
del test_embeddings
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

```
tensor([[ 0.3825,  0.1482,  0.6060, -0.5153,  0.4399,  0.0611, -0.6272, -0.0254,
          0.1643, -0.2210,  0.1442, -0.3721, -0.2168, -0.0890,  0.0979,  0.6561,
          0.6446,  0.4770,  0.8385,  1.6486,  0.8892, -0.1181, -0.0125, -0.5208,
          0.7785,  0.4872, -0.0150, -0.1413, -0.3475, -0.2959,  0.1028,  0.5719,
         -0.0456,  0.0264,  0.5382,  0.3226,  0.4079, -0.0436, -0.1460, -0.4835,
          0.3204,  0.5509, -0.7626,  0.4327,  0.6175, -0.3650, -0.6060, -0.7962,
          0.3929, -0.2367, -0.3472, -0.6120,  0.5475,  0.9481,  0.2094, -2.7771,
```

```
-0.6022, 0.8495, 1.2549, 0.0179, -0.0419, 2.1147, -0.0266, -0.2810,
0.6812, -0.1417, 0.9925, 0.4988, -0.6754, 0.6417, 0.4230, -0.2791,
0.0634, 0.6891, -0.3618, 0.0537, -0.1681, 0.1942, -0.4707, -0.1480,
-0.5899, -0.2797, 0.1679, 0.1057, -1.7601, 0.0088, -0.8333, -0.5836,
-0.3708, -0.5659, 0.2070, 0.0713, 0.0556, -0.2976, -0.0727, -0.2560,
0.4269, 0.0589, 0.0911, 0.4728]], grad_fn=<EmbeddingBackward0>)
```

```
[ ]: emmbed_dict = {}
number_lines = 0
with open("data/word-embeddings.txt") as f:
    for line in f:
        values = line.split()
        word = values[0]
        vector = np.asarray(values[1:], 'float32')
        emmbed_dict[word]=vector
        number_lines += 1
print(f"Number of words in Glove Embedding: {number_lines}")
```

Number of words in Glove Embedding: 29841

```
[ ]: len(emmbed_dict['in'])
# each word is associated with a vector of length 100
```

```
[ ]: 100
```

length of word ids is 32363 and the number of words in the embedding file is 29841; every word from the embedding file is used in the review dataset

2.5 2d Train and evaluate with pretraining

```
[ ]: # Fit and evaluate a model with pretrained embeddings without fine-tuning
class EmbeddingLSTM(nn.Module):
    """
    The RNN model that will be used to perform sentiment analysis.
    """

    def __init__(
        self,
        vocab_size,
        embedding_dim,
        hidden_dim,
        num_layers=1,
        lstm_dropout_prob=0.5,
        dropout_prob=0.3,
        finetuning=False
    ):
        """
        Initialize the model by setting up the layers
```

```

Parameters
-----
vocab_size: number of unique words in the reviews
embeddings_dim: size of the embeddings
hidden_dim: dimension of the LSTM output
num_layers: number of LSTM layers
lstm_dropout_prob: dropout applied between the LSTM layers
dropout_prob: dropout applied before the fully connected layer
"""

super().__init__()
self.finetuning = finetuning
self.num_layers = num_layers
self.hidden_dim = hidden_dim

self.embedding = nn.Embedding(
    num_embeddings=vocab_size,
    embedding_dim=embedding_dim
)

reviews_load_embeddings(embedding_layer=self.embedding,
↳word_ids=word_ids)

self.lstm = nn.LSTM(
    input_size=embedding_dim,
    hidden_size=hidden_dim,
    num_layers=num_layers,
    batch_first=True,
    dropout=lstm_dropout_prob
)

self.dropout = nn.Dropout(
    p=dropout_prob,
    inplace=False
)

self.fc = nn.Linear(
    in_features=hidden_dim,
    out_features=1
)

self.sigmoid = nn.Sigmoid()

# YOUR CODE HERE

def forward(self, x):

```

```

"""
Perform a forward pass of our model on some input and hidden state.

Parameters
-----
x: batch as a (batch_size, sequence_length) tensor

Returns
-----
Probability of positive class.
"""
# init hidden layer, which is needed for the LSTM
batch_size = len(x)
hidden = self.init_hidden(batch_size)

# YOUR CODE HERE

self.embedding.weight.requires_grad = self.finetuning
embedding = self.embedding(x)
out, (hidden, cell_state) = self.lstm(embedding, hidden)
out = self.dropout(out[:, -1, :])
out = self.fc(out)
out = self.sigmoid(out)
return out

def init_hidden(self, batch_size):
    """
    Initialize hidden state.

    Returns
    -----
    Empty hidden LSTM state.
    """

    # Create two new tensors with sizes num_layers x batch_size x
    ↪ hidden_dim,
    # initialized to zero, for hidden state and cell state of LSTM
    weight = next(self.parameters()) # only used to determine device

    hidden = (
        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_(),
        weight.new(self.num_layers, batch_size, self.hidden_dim).zero_(),
    )

    return hidden

```



```
[ ]: # Fit and evaluate a model with pretrained embeddings without fine-tuning.
# YOUR CODE HERE
vocab_size = len(word_ids) + 1 # +1 for the 0 padding
embedding_dim = 100
hidden_dim = 64
num_layers = 1

embedding_model = EmbeddingLSTM(vocab_size, embedding_dim, hidden_dim,
    ↪num_layers, finetuning=False).to(DEVICE)
reviews_train(embedding_model, train_loader, valid_loader, epochs=5,
    ↪device=DEVICE)
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Starting epoch 1

Epoch: 1/ 5	Batch: 5	Batch loss: 0.686438	Val loss: 0.706561	Val acc: 0.515000
Epoch: 1/ 5	Batch: 10	Batch loss: 0.695074	Val loss: 0.686327	Val acc: 0.527500
Epoch: 1/ 5	Batch: 15	Batch loss: 0.666638	Val loss: 0.684639	Val acc: 0.532500
Epoch: 1/ 5	Batch: 20	Batch loss: 0.679411	Val loss: 0.683167	Val acc: 0.547500
Epoch: 1/ 5	Batch: 25	Batch loss: 0.628898	Val loss: 0.688541	Val acc: 0.545000
Epoch: 1/ 5	Batch: 30	Batch loss: 0.701615	Val loss: 0.675063	Val acc: 0.570000
Epoch: 1/ 5	Batch: 35	Batch loss: 0.747103	Val loss: 0.689529	Val acc: 0.550000
Epoch: 1/ 5	Batch: 40	Batch loss: 0.657241	Val loss: 0.673139	Val acc: 0.570000
Epoch: 1/ 5	Batch: 45	Batch loss: 0.644778	Val loss: 0.669597	Val acc: 0.582500
Epoch: 1/ 5	Batch: 50	Batch loss: 0.652246	Val loss: 0.678360	Val acc: 0.597500
Epoch: 1/ 5	Batch: 55	Batch loss: 0.661263	Val loss: 0.719608	Val acc: 0.540000
Epoch: 1/ 5	Batch: 60	Batch loss: 0.702641	Val loss: 0.664131	Val acc: 0.572500

96

Finished epoch 1. Average batch loss: 0.6720854407176375. Average validation loss: 0.6848886503527561

Starting epoch 2

Epoch: 2/ 5	Batch: 65	Batch loss: 0.678968	Val loss: 0.684704	Val acc: 0.590000
Epoch: 2/ 5	Batch: 70	Batch loss: 0.598958	Val loss: 0.741897	Val acc: 0.520000

Epoch: 2/ 5	Batch: 75	Batch loss: 0.743488	Val loss: 0.681598	Val acc:
0.607500				
Epoch: 2/ 5	Batch: 80	Batch loss: 0.607069	Val loss: 0.657020	Val acc:
0.590000				
Epoch: 2/ 5	Batch: 85	Batch loss: 0.657863	Val loss: 0.651997	Val acc:
0.602500				
Epoch: 2/ 5	Batch: 90	Batch loss: 0.731028	Val loss: 0.658597	Val acc:
0.587500				
Epoch: 2/ 5	Batch: 95	Batch loss: 0.577248	Val loss: 0.663950	Val acc:
0.612500				
Epoch: 2/ 5	Batch: 100	Batch loss: 0.662151	Val loss: 0.715565	Val acc:
0.520000				
Epoch: 2/ 5	Batch: 105	Batch loss: 0.610220	Val loss: 0.663269	Val acc:
0.650000				
Epoch: 2/ 5	Batch: 110	Batch loss: 0.628149	Val loss: 0.650179	Val acc:
0.635000				
Epoch: 2/ 5	Batch: 115	Batch loss: 0.596552	Val loss: 0.646293	Val acc:
0.610000				
Epoch: 2/ 5	Batch: 120	Batch loss: 0.539033	Val loss: 0.621515	Val acc:
0.665000				
Epoch: 2/ 5	Batch: 125	Batch loss: 0.708193	Val loss: 0.645570	Val acc:
0.637500				

104

Finished epoch 2. Average batch loss: 0.6357951182872057. Average validation loss: 0.667857927771715

Starting epoch 3

Epoch: 3/ 5	Batch: 130	Batch loss: 0.581725	Val loss: 0.655778	Val acc:
0.625000				
Epoch: 3/ 5	Batch: 135	Batch loss: 0.534587	Val loss: 0.664557	Val acc:
0.612500				
Epoch: 3/ 5	Batch: 140	Batch loss: 0.614694	Val loss: 0.673386	Val acc:
0.595000				
Epoch: 3/ 5	Batch: 145	Batch loss: 0.570041	Val loss: 0.653104	Val acc:
0.625000				
Epoch: 3/ 5	Batch: 150	Batch loss: 0.675055	Val loss: 0.649794	Val acc:
0.622500				
Epoch: 3/ 5	Batch: 155	Batch loss: 0.597248	Val loss: 0.634350	Val acc:
0.642500				
Epoch: 3/ 5	Batch: 160	Batch loss: 0.583497	Val loss: 0.630251	Val acc:
0.657500				
Epoch: 3/ 5	Batch: 165	Batch loss: 0.585778	Val loss: 0.632745	Val acc:
0.657500				
Epoch: 3/ 5	Batch: 170	Batch loss: 0.667148	Val loss: 0.627554	Val acc:
0.657500				
Epoch: 3/ 5	Batch: 175	Batch loss: 0.560197	Val loss: 0.626309	Val acc:
0.662500				
Epoch: 3/ 5	Batch: 180	Batch loss: 0.580412	Val loss: 0.639708	Val acc:
0.642500				

Epoch: 3/ 5 Batch: 185 Batch loss: 0.577425 Val loss: 0.625761 Val acc: 0.665000
Epoch: 3/ 5 Batch: 190 Batch loss: 0.779278 Val loss: 0.649485 Val acc: 0.625000
104
Finished epoch 3. Average batch loss: 0.5963751999661326. Average validation loss: 0.6432909939724666
Starting epoch 4
Epoch: 4/ 5 Batch: 195 Batch loss: 0.526449 Val loss: 0.640628 Val acc: 0.637500
Epoch: 4/ 5 Batch: 200 Batch loss: 0.608472 Val loss: 0.644570 Val acc: 0.632500
Epoch: 4/ 5 Batch: 205 Batch loss: 0.664701 Val loss: 0.649488 Val acc: 0.622500
Epoch: 4/ 5 Batch: 210 Batch loss: 0.527539 Val loss: 0.618145 Val acc: 0.680000
Epoch: 4/ 5 Batch: 215 Batch loss: 0.519513 Val loss: 0.619536 Val acc: 0.675000
Epoch: 4/ 5 Batch: 220 Batch loss: 0.588045 Val loss: 0.626369 Val acc: 0.657500
Epoch: 4/ 5 Batch: 225 Batch loss: 0.554402 Val loss: 0.618975 Val acc: 0.672500
Epoch: 4/ 5 Batch: 230 Batch loss: 0.639828 Val loss: 0.627784 Val acc: 0.657500
Epoch: 4/ 5 Batch: 235 Batch loss: 0.578442 Val loss: 0.611042 Val acc: 0.680000
Epoch: 4/ 5 Batch: 240 Batch loss: 0.552774 Val loss: 0.628441 Val acc: 0.660000
Epoch: 4/ 5 Batch: 245 Batch loss: 0.494681 Val loss: 0.634671 Val acc: 0.632500
Epoch: 4/ 5 Batch: 250 Batch loss: 0.656134 Val loss: 0.699023 Val acc: 0.590000
Epoch: 4/ 5 Batch: 255 Batch loss: 0.574947 Val loss: 0.610381 Val acc: 0.675000
104
Finished epoch 4. Average batch loss: 0.5753812100738287. Average validation loss: 0.633004194841935
Starting epoch 5
Epoch: 5/ 5 Batch: 260 Batch loss: 0.566240 Val loss: 0.615236 Val acc: 0.687500
Epoch: 5/ 5 Batch: 265 Batch loss: 0.482931 Val loss: 0.609328 Val acc: 0.690000
Epoch: 5/ 5 Batch: 270 Batch loss: 0.456848 Val loss: 0.610621 Val acc: 0.667500
Epoch: 5/ 5 Batch: 275 Batch loss: 0.505860 Val loss: 0.605004 Val acc: 0.682500
Epoch: 5/ 5 Batch: 280 Batch loss: 0.534595 Val loss: 0.607091 Val acc: 0.677500

```

Epoch: 5/ 5    Batch: 285    Batch loss: 0.502120    Val loss: 0.613821 Val acc:
0.682500
Epoch: 5/ 5    Batch: 290    Batch loss: 0.599801    Val loss: 0.617692 Val acc:
0.660000
Epoch: 5/ 5    Batch: 295    Batch loss: 0.557508    Val loss: 0.614699 Val acc:
0.650000
Epoch: 5/ 5    Batch: 300    Batch loss: 0.545653    Val loss: 0.613553 Val acc:
0.677500
Epoch: 5/ 5    Batch: 305    Batch loss: 0.559772    Val loss: 0.665818 Val acc:
0.655000
Epoch: 5/ 5    Batch: 310    Batch loss: 0.538530    Val loss: 0.603941 Val acc:
0.672500
Epoch: 5/ 5    Batch: 315    Batch loss: 0.640959    Val loss: 0.601522 Val acc:
0.677500
Epoch: 5/ 5    Batch: 320    Batch loss: 0.585376    Val loss: 0.604607 Val acc:
0.670000
104
Finished epoch 5. Average batch loss: 0.5449454039335251. Average validation
loss: 0.6140717204946738

```

```
[ ]: reviews_test(embedding_model, test_loader, device=DEVICE)
```

```

Test loss: 0.548
Test accuracy: 0.725

```

```

[ ]: # Fit and evaluate a model with pretrained embeddings with fine-tuning.
# YOUR CODE HERE
vocab_size = len(word_ids) + 1 # +1 for the 0 padding
embedding_dim = 100
hidden_dim = 64
num_layers = 1

embedding_model = EmbeddingLSTM(vocab_size, embedding_dim, hidden_dim,
    ↪ num_layers, finetuning=True).to(DEVICE)
reviews_train(embedding_model, train_loader, valid_loader, epochs=5,
    ↪ device=DEVICE)

```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Starting epoch 1

```

Epoch: 1/ 5    Batch: 5    Batch loss: 0.675380    Val loss: 0.703210 Val acc:
0.502500
Epoch: 1/ 5    Batch: 10   Batch loss: 0.673556    Val loss: 0.688418 Val acc:
0.552500
Epoch: 1/ 5    Batch: 15   Batch loss: 0.688058    Val loss: 0.686013 Val acc:
0.530000
Epoch: 1/ 5    Batch: 20   Batch loss: 0.692807    Val loss: 0.685314 Val acc:
0.555000

```

Epoch: 1/ 5	Batch: 25	Batch loss: 0.675730	Val loss: 0.689419	Val acc: 0.520000
Epoch: 1/ 5	Batch: 30	Batch loss: 0.663634	Val loss: 0.682638	Val acc: 0.555000
Epoch: 1/ 5	Batch: 35	Batch loss: 0.671265	Val loss: 0.676966	Val acc: 0.600000
Epoch: 1/ 5	Batch: 40	Batch loss: 0.675061	Val loss: 0.672660	Val acc: 0.602500
Epoch: 1/ 5	Batch: 45	Batch loss: 0.633352	Val loss: 0.673495	Val acc: 0.580000
Epoch: 1/ 5	Batch: 50	Batch loss: 0.666285	Val loss: 0.668230	Val acc: 0.622500
Epoch: 1/ 5	Batch: 55	Batch loss: 0.777652	Val loss: 0.661762	Val acc: 0.597500
Epoch: 1/ 5	Batch: 60	Batch loss: 0.606284	Val loss: 0.665211	Val acc: 0.580000

96

Finished epoch 1. Average batch loss: 0.6720362938940525. Average validation loss: 0.6794445620228847

Starting epoch 2

Epoch: 2/ 5	Batch: 65	Batch loss: 0.541709	Val loss: 0.669934	Val acc: 0.585000
Epoch: 2/ 5	Batch: 70	Batch loss: 0.615435	Val loss: 0.656927	Val acc: 0.595000
Epoch: 2/ 5	Batch: 75	Batch loss: 0.510195	Val loss: 0.638432	Val acc: 0.647500
Epoch: 2/ 5	Batch: 80	Batch loss: 0.696451	Val loss: 0.651840	Val acc: 0.597500
Epoch: 2/ 5	Batch: 85	Batch loss: 0.509652	Val loss: 0.681286	Val acc: 0.592500
Epoch: 2/ 5	Batch: 90	Batch loss: 0.581267	Val loss: 0.771913	Val acc: 0.552500
Epoch: 2/ 5	Batch: 95	Batch loss: 0.661655	Val loss: 0.630278	Val acc: 0.650000
Epoch: 2/ 5	Batch: 100	Batch loss: 0.585351	Val loss: 0.619743	Val acc: 0.667500
Epoch: 2/ 5	Batch: 105	Batch loss: 0.496805	Val loss: 0.747108	Val acc: 0.590000
Epoch: 2/ 5	Batch: 110	Batch loss: 0.524939	Val loss: 0.636721	Val acc: 0.655000
Epoch: 2/ 5	Batch: 115	Batch loss: 0.621672	Val loss: 0.610787	Val acc: 0.680000
Epoch: 2/ 5	Batch: 120	Batch loss: 0.434247	Val loss: 0.586571	Val acc: 0.695000
Epoch: 2/ 5	Batch: 125	Batch loss: 0.607140	Val loss: 0.586273	Val acc: 0.687500

104

Finished epoch 2. Average batch loss: 0.5507332952693105. Average validation

loss: 0.6529088149277064

Starting epoch 3

Epoch: 3/ 5	Batch: 130	Batch loss: 0.453490	Val loss: 0.679404	Val acc: 0.617500
Epoch: 3/ 5	Batch: 135	Batch loss: 0.583140	Val loss: 0.575940	Val acc: 0.717500
Epoch: 3/ 5	Batch: 140	Batch loss: 0.358122	Val loss: 0.831501	Val acc: 0.580000
Epoch: 3/ 5	Batch: 145	Batch loss: 0.411806	Val loss: 0.582739	Val acc: 0.690000
Epoch: 3/ 5	Batch: 150	Batch loss: 0.384330	Val loss: 0.655523	Val acc: 0.705000
Epoch: 3/ 5	Batch: 155	Batch loss: 0.344274	Val loss: 0.620297	Val acc: 0.695000
Epoch: 3/ 5	Batch: 160	Batch loss: 0.358667	Val loss: 0.822217	Val acc: 0.625000
Epoch: 3/ 5	Batch: 165	Batch loss: 0.310987	Val loss: 0.616756	Val acc: 0.727500
Epoch: 3/ 5	Batch: 170	Batch loss: 0.349381	Val loss: 0.570314	Val acc: 0.717500
Epoch: 3/ 5	Batch: 175	Batch loss: 0.404174	Val loss: 0.572764	Val acc: 0.732500
Epoch: 3/ 5	Batch: 180	Batch loss: 0.485828	Val loss: 0.653823	Val acc: 0.662500
Epoch: 3/ 5	Batch: 185	Batch loss: 0.266039	Val loss: 0.603032	Val acc: 0.732500
Epoch: 3/ 5	Batch: 190	Batch loss: 0.398966	Val loss: 0.558743	Val acc: 0.730000

104

Finished epoch 3. Average batch loss: 0.38209939328953624. Average validation

loss: 0.6417732508136675

Starting epoch 4

Epoch: 4/ 5	Batch: 195	Batch loss: 0.220744	Val loss: 0.575362	Val acc: 0.745000
Epoch: 4/ 5	Batch: 200	Batch loss: 0.275859	Val loss: 0.575091	Val acc: 0.737500
Epoch: 4/ 5	Batch: 205	Batch loss: 0.330835	Val loss: 0.572881	Val acc: 0.737500
Epoch: 4/ 5	Batch: 210	Batch loss: 0.106531	Val loss: 0.623701	Val acc: 0.737500
Epoch: 4/ 5	Batch: 215	Batch loss: 0.256790	Val loss: 0.640666	Val acc: 0.672500
Epoch: 4/ 5	Batch: 220	Batch loss: 0.637759	Val loss: 0.679671	Val acc: 0.735000
Epoch: 4/ 5	Batch: 225	Batch loss: 0.208216	Val loss: 0.641050	Val acc: 0.747500
Epoch: 4/ 5	Batch: 230	Batch loss: 0.265210	Val loss: 0.604293	Val acc: 0.735000

Epoch: 4/ 5 Batch: 235 Batch loss: 0.372622 Val loss: 0.641315 Val acc: 0.737500
 Epoch: 4/ 5 Batch: 240 Batch loss: 0.205042 Val loss: 0.677802 Val acc: 0.670000
 Epoch: 4/ 5 Batch: 245 Batch loss: 0.379012 Val loss: 0.648437 Val acc: 0.752500
 Epoch: 4/ 5 Batch: 250 Batch loss: 0.270986 Val loss: 0.580526 Val acc: 0.742500
 Epoch: 4/ 5 Batch: 255 Batch loss: 0.232753 Val loss: 0.571961 Val acc: 0.747500

104

Finished epoch 4. Average batch loss: 0.27838339447043836. Average validation loss: 0.6179042008633797

Starting epoch 5

Epoch: 5/ 5 Batch: 260 Batch loss: 0.209052 Val loss: 0.580086 Val acc: 0.752500
 Epoch: 5/ 5 Batch: 265 Batch loss: 0.161226 Val loss: 0.627696 Val acc: 0.730000
 Epoch: 5/ 5 Batch: 270 Batch loss: 0.161622 Val loss: 0.694488 Val acc: 0.727500
 Epoch: 5/ 5 Batch: 275 Batch loss: 0.195214 Val loss: 0.645984 Val acc: 0.742500
 Epoch: 5/ 5 Batch: 280 Batch loss: 0.227829 Val loss: 0.684968 Val acc: 0.742500
 Epoch: 5/ 5 Batch: 285 Batch loss: 0.155799 Val loss: 0.556669 Val acc: 0.765000
 Epoch: 5/ 5 Batch: 290 Batch loss: 0.174857 Val loss: 0.586511 Val acc: 0.732500
 Epoch: 5/ 5 Batch: 295 Batch loss: 0.158507 Val loss: 0.726699 Val acc: 0.702500
 Epoch: 5/ 5 Batch: 300 Batch loss: 0.325574 Val loss: 0.714649 Val acc: 0.695000
 Epoch: 5/ 5 Batch: 305 Batch loss: 0.169596 Val loss: 0.574627 Val acc: 0.755000
 Epoch: 5/ 5 Batch: 310 Batch loss: 0.354525 Val loss: 0.598333 Val acc: 0.727500
 Epoch: 5/ 5 Batch: 315 Batch loss: 0.179471 Val loss: 0.694602 Val acc: 0.692500
 Epoch: 5/ 5 Batch: 320 Batch loss: 0.142213 Val loss: 0.628788 Val acc: 0.732500

104

Finished epoch 5. Average batch loss: 0.19638578174635768. Average validation loss: 0.6395460776984692

```
[ ]: reviews_test(embedding_model, test_loader, device=DEVICE)
```

Test loss: 0.558

Test accuracy: 0.770

2.6 2e Sandbox

```
[ ]: # Explore different architectures and hyperparameters.
```